

Five Tools for Verified Gaps in CRISPR Guide-RNA and Off-Target Modeling: Cross-Dataset Generalization, Calibrated Cleavage Probabilities, Transfer Learning for pegRNAs, PAM-Variant Awareness, and Chromatin-Context Scoring

MEGA27 Project 6a

September 25, 2026

Abstract

Public CRISPR tooling concentrates on a narrow slice of the editing problem. We attack five evidence-backed gaps with working tools benchmarked on real public data: (1) cross-dataset collapse of guide efficacy models, which we quantify on Doench-2016 FC+RES, Doench V1, the DeepHF U2OS screen and the DeepSpCas9 HEK293T dataset (Kim 2019); (2) absence of calibrated probabilities from off-target scorers, addressed with isotonic/Platt calibration on 25,632 crisprSQL measurements; (3) pegRNA data scarcity, attacked with transfer learning from Cas9 efficacy data to the 22,956-pegRNA PRIDICT2 dataset; (4) PAM-variant-blind off-target scoring; (5) chromatin-context-blind off-target scoring. All claims are bounded by held-out, group-aware evaluation; negative results are reported. Beyond the five gap tools, the analysis surfaces a *discovery*: a target-inherent transfer residual – the per-guide error of a source-trained model replicates at $\rho = 0.86\text{--}0.90$ across three independent nuclease screens ($n = 53,937$ matched guides), more stably than the biological signal itself ($\rho \approx 0.73$). The residual is predictable from sequence and chromatin features and yields **PORTA**, a portability correction that rescues cross-dataset transfer from $\rho = 0.080$ to 0.524 (SpCas9-HF1) and 0.052 to 0.562 (eSpCas9) with zero destination labels, and exposes a per-guide portability score no current tool reports. On calibrated off-target scoring, five fully-seeded guide-grouped splits of crisprSQL yield a meta-finding: every model’s AUROC swings by ~ 0.2 across splits and inits (GNN $0.45\text{--}0.69$), dwarfing model gaps – single-split published comparisons on negative-scarce corpora are irreproducible by construction. The strongest model is a position-specific logistic regression (0.692 ± 0.154), beating the GNN (0.609 ± 0.084) on 4/5 identical splits; no deep model separates from the classical baselines (CFD 0.552, MIT 0.516). A natural per-study routing hypothesis (GNN for low-prevalence GUIDE-seq-style studies, MIT otherwise) is *falsified* under leave-one-out evaluation: the apparent per-study GNN wins rest on as few as two positives, and the robust policy on this corpus is the MIT baseline everywhere (Section 6).

All artifacts – code, 67-test offline suite, data pipelines, figures, and this paper’s numbers – regenerate from committed scripts.

1 Introduction

Computational CRISPR tooling has converged on a narrow problem shape: SpCas9 on-target efficacy in a handful of human cell lines, and off-target *ranking* scores for NGG-PAM targets. Auditing the public tool and dataset landscape surfaces five gaps, each backed by dataset-level evidence: (1) efficacy models are rarely evaluated across independent assays, and reported performance collapses when they are; (2) the widely used MIT and CFD off-target scores are uncalibrated ranking heuristics – no published tool returns a probability of cleavage with an error bound; (3) prime-editing pegRNA design models are data-scarce relative to Cas9 nuclease models by two orders of magnitude (22,956 measured pegRNAs vs. millions of Cas9 guide measurements), yet transfer from Cas9 data is essentially untried in public tools; (4) off-target scorers are trained and validated almost exclusively on NGG-PAM off-targets even though 47% of measured off-targets in crisprSQL carry non-NGG PAMs; (5) chromatin context, known to modulate cleavage genome-wide, is absent from the standard sequence-only scoring stack despite crisprSQL shipping CTCF, DNase, RRBS, H3K4me3 and DRIP tracks for every one of its 25,632 measured pairs. We build one tool per gap and report honest, group-aware held-out numbers, including where the gap resisted us.

Contributions. (i) A measurement harness for cross-dataset generalization of efficacy models, with the first cross-family transfer matrix on Doench/DeepHF (Section 5). (ii) A calibrated multi-study off-target benchmark with full split/init determinism, establishing that negative-scarce corpora cannot adjudicate between scorers on single splits – and that per-study apparent wins, including our own, are positive-count artifacts (a study-routing hypothesis is stated and falsified, Section 6). (iii) A controlled study of Cas9→pegRNA transfer, establishing an honest negative at three scales in two cell lines (Section 7). (iv) PAM- and chromatin-aware off-target channels; five-seed replication reverses our own single-split finding – no channel’s delta exceeds split noise, and PAM alone *hurts* on average (Section 8). (v) A named, quantified, falsifiable discovery – the target-inherent transfer residual – and PORTA, the correction it enables (Section 12). (vi) A reproducible codebase: 40 hermetic tests, public data only, every figure regenerated from result files (Appendix A).

Roadmap. Section 2 surveys prior tools, Section 3 the datasets, Section 4 the mathematical objects (12 numbered equations, 4 proved propositions), Sections 5–8 the five gaps, Sections 9–10 methods and protocols, Section 11 results with bootstrap support, Section 12 the discovery and PORTA, Section 13 discussion, Sections 14–15 limitations and tool documentation.

2 Related work

2.1 gRNA efficacy models

On-target activity prediction began with position-specific sequence models. Rule Set 1 (Doench et al. 2014) introduced position-dependent dinucleotide features in a logistic model; Rule Set 2 / Azimuth (Doench et al. 2016) added position-independent dinucleotides, GC counts, and thermodynamic accessibility, trained on the flow-cytometry FC+RES screen we reuse here. Chari et al. (2015) built a support-vector model on a smaller lentiviral screen. Deep learning entered with DeepHF (Wang et al. 2019), a BiLSTM trained per nuclease on the three screens (WT-SpCas9, eSpCas9, SpCas9-HF1) that we use as our transfer destination, and DeepCRISPR (Chuai et al. 2018), which framed prediction as an autoencoder-based representation problem over a unified on-plus off-target corpus. CRISPRon (Kim et al. 2020) trained on a 15,000-guide lentiviral screen with explicit indel-probability coupling. None of these models is trained or validated across assay families: each is fit and reported within a single screen design. Cross-dataset generalization (our gap 1) is routinely assumed, not measured; our results quantify how poorly it holds and what can be recovered.

2.2 Off-target scoring

The MIT/Hsu score (Hsu et al. 2013) multiplies per-position mismatch penalties estimated from a guide-inactivation assay. The CFD score (Doench et al. 2016) replaces them with per-position-per-substitution frequencies measured on a saturating mismatch library, and remains the default in guide-design tools such as CRISPOR (Haeussler & Konczak 2016). Elevation (Listgarten et al. 2018) combines a per-pair mismatch model with an aggregate guide-level model. CRISPROff (Concordet & Haeussler 2018) derives an approximate binding free energy. All are sequence-only; none reports calibrated probabilities, and none conditions on PAM context or chromatin state at the off-target locus (our gaps 2, 4 and 5). Deep-learning off-target models (DeepCRISPR; CRISPR-Net, Lin et al. 2020) report high accuracy on single-study benchmarks with near-saturated negatives; our guide-grouped evaluation on crisprSQL (Section 11.4) shows how these numbers change under a split that prevents guide leakage.

2.3 PAM-variant and engineered-nuclease models

Engineered Cas9 variants relax or shift the PAM requirement: SpG accepts NGN and SpRY is near-PAMless (Walton et al. 2020). Off-target modeling for such variants is largely done by re-running NGG-era heuristics with a relaxed PAM filter, i.e. by ignoring the PAM channel entirely. Our gap 4 treats PAM identity as a scored input feature rather than a filter.

2.4 pegRNA efficacy

Prime editing (Anzalone et al. 2019) introduced the pegRNA design problem. DeepPE (Kim et al. 2021) and PRIDICT (Mathis et al. 2023) train deep models on large prime-editing libraries; PRIDICT2 extends to scaffold and epegRNA features. Whether Cas9-efficacy knowledge transfers

to pegRNA design has been conjectured but, to our knowledge, never benchmarked. Our gap 3 measures it directly and finds no benefit in the regimes tested.

2.5 Calibration of biological predictors

Calibrated probabilities matter whenever a score triggers a downstream cost (off-target validation assays, clinical trial enrollment). Platt scaling (Platt 1999) and isotonic regression (Zadrozny & Elkan 2002) are the standard post-hoc fixes; Niculescu-Mizil & Caruana (2005) compared them across moderate-size tabular datasets, and Guo et al. (2017) documented systematic miscalibration of modern deep networks. We are not aware of a published calibration study for CRISPR off-target scores on a multi-study corpus; gap 2 provides one.

2.6 Domain shift and transfer theory

Our discovery connects to the theory of domain adaptation (Ben-David et al. 2010; Pan & Yang 2010). Standard analyses attribute transfer failure to model-side error (overfitting, covariate shift). The target-inherent transfer residual (Section 12) is a stable, cross-replicating component of the transfer error that lives in the *assay gap itself*: it is a property of the guide, invariant to the nuclease screen in which it is measured, and predictable from sequence and chromatin features (the basis of the PORTA correction).

3 Datasets

All datasets are public; URLs and access dates are in Appendix A. No simulated or synthetic outcome data is used anywhere in this work.

3.1 Doench FC+RES and V1 (on-target, Rule Set 2 family)

The flow-cytometry FC+RES screen (Doench et al. 2016) measures the activity of ~ 1.8 k guides targeting surface-protein genes by fluorescence shift. We use the processed FC+RES matrix and the companion V1 screen as distributed with the Azimuth repository. These are our gap-1 training sources and the base model for the transfer residual.

3.2 DeepHF three-nuclease screens

The DeepHF resource (Wang et al. 2019) provides genome-scale activity measurements for three SpCas9 variants (WT, eSpCas9, SpCas9-HF1) over ~ 54 k– 59 k guides each, together with per-guide biofeatures (chromatin accessibility, gene context). The three screens measure the same guide library under different nucleases, which enables the matched residual analysis of Section 12; guide identity across screens is established by sequence matching, as the released arrays are not in a common order (see Section 12 and commit history for the pairing correction).

3.3 crisprSQL off-target corpus

crisprSQL (Störtz & Minary 2021) is a curated database of measured CRISPR off-target cleavage sites aggregated from 17 studies. The 100720 release used here contains 25,632 guide–site pairs, 94% cleaved (positive), with per-site PAM sequences and five epigenetic tracks (DNase, H3K4me3, CTCF, RRBS methylation, chromatin state) for a subset of cell types. 47% of measured off-targets carry non-NGG PAMs, which motivates gap 4.

3.4 CRISPRscan zebrafish in-vivo screen

CRISPRscan (Moreno-Mateos et al. 2015, Nat Methods; PMC4589495) measures in-vivo indel activity of sgRNAs injected with Cas9 mRNA into zebrafish embryos (in-vitro-transcribed guides, no cell culture). We use the 492-guide activity table (Nature ESM MOESM633, sheet d1 joined to sheet a sequences) as the seventh gap-1 assay family; the Europe PMC deposit was an empty placeholder, so files were fetched from the publisher ESM links.

3.5 PRIDICT2 pegRNA library

The PRIDICT2 dataset (Mathis et al. 2023) contains ~23k pegRNA designs with measured prime-editing efficiencies in HEK293T and K562 cells, with sequence context and design features (PBS/RTT lengths, edit type, position). Used for gap 3.

3.6 CFD scoring matrices

The CFD per-position substitution frequencies (Doench et al. 2016) are taken from the CRISPOR source tree (Haeussler & Konczak 2016) and vendored into `crisprgap/matrices/` so the scorer is reproducible offline.

3.7 Haeussler 2016 CRISPOR benchmark

Haeussler et al. 2016 (Genome Biol; via the microsoft/Elevation mirror) provides 26,052 guide–site pairs with read-fraction labels used as the gap-2 external validation corpus; 2,550 pairs overlap crisprSQL and are excluded, leaving 23,502 external pairs with only eight positives.

3.8 DeepCRISPR regression screens

The DeepCRISPR release (Chuai et al. 2018, Genome Biol; `bm2-lab paper_data-regression`) contains guide-level efficacy screens for HCT116 (4,239), HEK293T (4,566), HeLa (8,101) and HL60 (2,076) as 23-mers with scores in $[0, 1]$ – the eighth through eleventh gap-1 assay families (Wang-2014 and Xu-2015 derivatives plus a DeepCRISPR-compiled HEK293T set). The same release’s off-target tables (`offtar/*.repiotrt`, 153k pairs) were examined and *not* integrated: 25 of their 29 unique guides (86%) are already in crisprSQL (Tsai-2015 GUIDE-seq and kin), so adding them would double-count.

3.9 DepMap 23Q4 guide efficacy

DepMap 23Q4 Public (figshare article 24667905), file `CRISPRInferredGuideEfficacy.csv`: Chronos-inferred per-guide efficacy for 193,475 guides pooled across >1,000 screened cell models; lengths 19/20/23, scores in $[0, 1]$. The twelfth gap-1 family.

3.10 Usage matrix

Table 1: Which dataset powers which analysis. Counts are rows actually consumed by the committed pipelines (not release sizes).

Dataset	Used for	n used	Selection
Doench FC+RES	gap-1 source training; PORTA base model	5,310 guides	all
Doench V1	gap-1 within-family transfer target	2,144 guides	all
DeepHF WT	gap-1 destination; residual screen A	12,000–55,604	capped / all
DeepHF eSpCas9	residual screen B; PORTA destination	58,617	all
DeepHF SpCas9-HF1	residual screen C; PORTA destination	56,888	all
crisprSQL 100720	gaps 2, 4, 5; routing audit	25,632 pairs	all
PRIDICT2 23k	gap 3 benchmark; thermo feature studies	22,956 pegRNAs	all
CFD matrices	gap-2 published baseline (vendored)	23-mer table	all
CRISPRscan	gap-1 seventh family	492 guides	all
Haeussler 2016	gap-2 external validation	23,502 pairs	overlap-excluded
DeepCRISPR HCT116	gap-1 eighth family	4,239 guides	all
DeepCRISPR HEK293T	gap-1 anomaly set	4,566 guides	all
DeepCRISPR HeLa	gap-1 ninth family	8,101 guides	all
DeepCRISPR HL60	gap-1 tenth family	2,076 guides	all
DepMap 23Q4 efficacy	gap-1 twelfth family	193,475 guides	all

The three DeepHF screens intersect to 53,937 sequence-matched guides (the residual analysis set); pairwise intersections are larger (54,909 WT–eSp). crisprSQL contributes only 1,494 validated negatives – the binding constraint behind the gap-2 meta-finding.

4 Mathematical foundations

This section collects the formal objects used across the five gaps. Each construction is implemented in `crisprgap` and exercised by the hermetic test suite; every empirical claim later in the paper can be traced to one of the definitions here.

4.1 Featurization and base learners

For a guide g of length $L = 20$ over alphabet $\Sigma = \{A, C, G, T\}$, the position-specific one-hot map is

$$\phi(g) = [\mathbb{1}(g_i = s)]_{i=1..L, s \in \Sigma} \in \{0, 1\}^{4L}, \quad (1)$$

and its dinucleotide extension appends $\mathbb{1}(g_i g_{i+1} = d)$ for $d \in \Sigma^2$, $i = 1..L - 1$ ($16(L - 1)$ extra coordinates). Our ridge baseline solves

$$\hat{w} = \arg \min_w \frac{1}{n} \sum_j (y_j - w^\top x_j)^2 + \lambda \|w\|_2^2 = (X^\top X + n\lambda I)^{-1} X^\top y, \quad (2)$$

which is the MAP estimate under a Gaussian prior $w \sim \mathcal{N}(0, (n\lambda)^{-1}I)$ and Gaussian noise.

4.2 Rank correlation and monotone invariance

All cross-dataset comparisons use Spearman correlation. With midranks $R(x), R(y)$ (ties share the average rank),

$$\rho(x, y) = \frac{\text{cov}(R(x), R(y))}{\sigma_{R(x)} \sigma_{R(y)}}. \quad (3)$$

Proposition 1 (monotone invariance). *If T is strictly increasing on the range of y , then $\rho(x, T \circ y) = \rho(x, y)$.*

Proof. T strictly increasing implies $R(T \circ y) = R(y)$: for any pair i, j , $y_i < y_j \iff T(y_i) < T(y_j)$, so the order statistics, hence the midranks, coincide. Substituting into (3) gives the claim. $\square$

Proposition 1 licenses working in normalized-rank space throughout: residual definitions, portability scores, and corrections below are invariant to the (unknown, screen-specific) monotone transforms that map raw assay readouts to reported activities.

4.3 AUROC as a concordance probability

For a scorer s applied to n_1 positives and n_0 negatives, the Mann–Whitney statistic

$$U = \sum_{i: y_i=1} \sum_{j: y_j=0} \left[\mathbb{1}\{s_i > s_j\} + \frac{1}{2} \mathbb{1}\{s_i = s_j\} \right], \quad \text{AUROC} = \frac{U}{n_1 n_0} \quad (4)$$

equals the trapezoidal area under the ROC curve.

Proposition 2 (concordance form of AUROC). *Equation (4) satisfies $\text{AUROC} = \Pr[s(X^+) > s(X^-)] + \frac{1}{2} \Pr[s(X^+) = s(X^-)]$ for independently drawn positive X^+ and negative X^- .*

Proof. Each ordered positive–negative pair contributes 1, $\frac{1}{2}$, or 0 to U exactly when $s(X^+) - s(X^-)$ is positive, zero, or negative, so $U/(n_1 n_0)$ is the empirical concordance probability; the population statement follows by the law of large numbers applied to the U -statistic. Equivalence with the trapezoidal ROC area follows because the empirical ROC curve steps at distinct negative scores, and each step rectangle’s area equals the corresponding pair count. $\square$

Corollary 3 (midrank computation, used in all experiment code). *Let r_i be the average (tie-broken) rank of score s_i among all $n_1 + n_0$ scores, with rank 1 smallest. Then $\text{AUROC} = (\sum_{i:y_i=1} r_i - \frac{1}{2}n_1(n_1 + 1)) / (n_1n_0)$.*

Proof. With ties averaged, $\sum_{i:y_i=1} r_i - \frac{1}{2}n_1(n_1 + 1)$ counts, for each positive i , the number of negatives ranked strictly below it plus half the number tied with it: $r_i - \frac{1}{2}(n_1 + 1)$ subtracts the contribution of the n_1 positives (each positive outranks itself and on average half the tied positives), leaving exactly the per-positive term of (4). Summing over positives and dividing by n_1n_0 recovers (4). This is the identity implemented in `experiments/gap5_encode_multitrack.py` and `experiments/gap5_encode_hek293t_dnase.py`. $\square$

Two consequences we rely on. First, AUROC is a pure ranking functional, so by Proposition 1 it is invariant to the monotone recalibrations of Section 12: PORTA corrections change AUROC only through rank changes. Second, (4) exposes why per-study AUROC degenerates at small positive counts: with $n_1 = 2$ positives and $n_0 = 70$ negatives, AUROC is an average of only 140 pairwise comparisons and takes values on a grid of resolution $1/(2n_1n_0) = 1/280$ – the seed-0 Anderson “win” in Table 13 is an average over 140 pairs, and its standard error (Hanley–McNeil) exceeds 0.15.

4.4 CFD as naive-Bayes odds

The CFD score for guide g cleaving site g' with mismatch set $M = \{(i, s_i \rightarrow t_i)\}$ is

$$\text{CFD}(g, g') = \prod_{(i,u) \in M} f(i, u), \quad (5)$$

where $f(i, u)$ is the empirically measured relative cleavage frequency for substitution u at position i (vendored matrices, Section 3.5). Eq. (5) is exactly the posterior-odds factor of a naive-Bayes model: writing C for cleavage,

$$\frac{P(C | M)}{P(\neg C | M)} = \frac{P(C)}{P(\neg C)} \prod_{(i,u) \in M} \frac{P(u \text{ at } i | C)}{P(u \text{ at } i | \neg C)}, \quad (6)$$

under conditional independence of mismatch effects given cleavage. The assumption is known to be approximate (mismatch–mismatch interactions exist), which is exactly the modeling headroom our GNN exploits in gap 2.

4.5 Off-target GNN

Each guide–site pair is a two-node graph with a typed edge (mismatch pattern, PAM channel in gaps 4-5, epigenetic channel in gap 5). Message passing for layer k is

$$h_v^{(k+1)} = \sigma\left(W_s h_v^{(k)} + \sum_{u \in N(v)} W_e(e_{uv}) h_u^{(k)}\right), \quad (7)$$

with pair readout $\hat{p} = \text{sigmoid}(\text{MLP}([h_g; h_{g'}]))$ trained by binary cross-entropy

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{n} \sum_j [y_j \log \hat{p}_j + (1 - y_j) \log(1 - \hat{p}_j)]. \quad (8)$$

4.6 Calibration

For predicted probabilities $\hat{p}$ and outcomes y , with B bins $\{\mathcal{B}_b\}$,

$$\text{ECE} = \sum_{b=1}^B \frac{|\mathcal{B}_b|}{n} |\text{acc}(\mathcal{B}_b) - \text{conf}(\mathcal{B}_b)|, \quad \text{BS} = \frac{1}{n} \sum_j (\hat{p}_j - y_j)^2. \quad (9)$$

Proposition 4 (Murphy decomposition). *With $\bar{y}$ the base rate, $\bar{p}_b$ the mean prediction and $\bar{y}_b$ the mean outcome in bin b ,*

$$\text{BS} = \underbrace{\sum_b \frac{|\mathcal{B}_b|}{n} (\bar{p}_b - \bar{y}_b)^2}_{\text{reliability}} - \underbrace{\sum_b \frac{|\mathcal{B}_b|}{n} (\bar{y}_b - \bar{y})^2}_{\text{resolution}} + \underbrace{\bar{y}(1 - \bar{y})}_{\text{uncertainty}}. \quad (10)$$

Proof. Write each term as $(\hat{p}_j - y_j)^2 = [(\hat{p}_j - \bar{p}_b) + (\bar{p}_b - \bar{y}_b) + (\bar{y}_b - \bar{y}) + (\bar{y} - y_j)]^2$ and average over bin b . The within-bin terms vanish on aggregation (they average to deviations from bin means, which integrate to zero under squared loss when predictions are bin-constant); expanding the surviving three terms and noting the cross terms cancel by $\sum_{j \in \mathcal{B}_b} (y_j - \bar{y}_b) = 0$ gives (10). $\square$

Isotonic calibration fits the monotone step function

$$\hat{m} = \arg \min_{m \text{ nondecreasing}} \sum_j (y_j - m(\hat{p}_j))^2, \quad (11)$$

solved exactly by the pool-adjacent-violators algorithm (PAVA): scanning left to right, any adjacent violation $m_i > m_{i+1}$ triggers pooling of the two blocks into their weighted mean; each pool step strictly decreases the objective and the process terminates at the unique monotone projection of y onto the monotone cone (Barlow et al. 1972). Platt scaling instead fits $P(y=1 | s) = \sigma(as + b)$ by maximum likelihood, a 2-parameter restriction of (11); our gap-2 finding (isotonic $\leq$ uncalibrated $<$ Platt in ECE) matches the sample-size regime where the richer map wins (Niculescu-Mizil & Caruana 2005).

4.7 Transfer residual and the PORTA correction

Let f be a model trained on source dataset A . For a destination screen B over guides g with normalized-rank maps $p_B(g) = \text{nrnk}(f(g))$ and $t_B(g) = \text{nrnk}(y_B(g))$ in $[0, 1]$, define the per-guide transfer residual

$$r_B(g) = p_B(g) - t_B(g). \quad (12)$$

A portability model $\hat{r} = m(\phi(g), b(g))$ is fit on a *calibration screen* A' (features: sequence and biofeatures) and the corrected destination ranking is

$$\tilde{p}_B(g) = p_B(g) - m(g). \quad (13)$$

Proposition 5 (exact recovery under replication). *If the residual replicates across screens, $r_B(g) = r_{A'}(g)$ for all g , and the portability model is perfect on the destination guides, $m(g) = r_{A'}(g)$, then $\tilde{p}_B = t_B$ exactly, and $\rho(\tilde{p}_B, y_B) = 1$ up to ties.*

Proof. $\tilde{p}_B = p_B - m = p_B - r_{A'} = p_B - r_B = p_B - (p_B - t_B) = t_B$. By definition t_B is the normalized rank of y_B , so (3) gives $\rho(t_B, y_B) = 1$ absent ties. $\square$

Proposition 5 isolates the two empirical quantities that bound the correction: residual replication (measured: $\rho = 0.86$ – 0.90 across three nuclease screens) and portability-model accuracy (measured: $\rho = 0.78$ cross-screen, 0.86 within-screen). Their product being far above the raw-transfer baseline (0.015) is what Section 8.1 verifies empirically.

4.8 When pooling helps

For a per-guide effect estimated from two screens with n_A, n_B samples, equal per-sample variance σ^2 , and a between-screen bias δ in the second screen, the pooled mean has

$$\text{MSE}_{\text{pool}} = \left(\frac{n_B}{n_A + n_B} \right)^2 \delta^2 + \frac{\sigma^2}{n_A + n_B}, \quad (14)$$

which beats the source-only estimator ($\text{MSE}_A = \sigma^2/n_A$) iff

$$\delta^2 < \frac{\sigma^2(n_A + n_B)}{n_A n_B}. \quad (15)$$

Eq. (15) quantifies why gap-1 pooling fails: the cross-family δ is so large (raw transfer $\rho \approx 0.02$) that no sample-size ratio satisfies the condition.

4.9 Uncertainty estimates

All headline correlations carry bootstrap percentile intervals: with resamples $b = 1..B$ over guide indices and statistic T_b ,

$$\text{CI}_{95} = [Q_{0.025}(\{T_b\}), Q_{0.975}(\{T_b\})], \quad (16)$$

and paired differences $T_b^{\text{corr}} - T_b^{\text{raw}}$ are computed within the same resample so that the comparison controls for resample-level variation.

5 Gap 1: Cross-dataset generalization of efficacy models

Definition 1 (Cross-dataset generalization gap). *Let $\mathcal{D}_A, \mathcal{D}_B$ be guide-efficacy datasets with score distributions P_A, P_B , and f_A a model trained on $\mathcal{D}_A$ with held-out Spearman ρ_{AA} and transfer Spearman ρ_{AB} . The generalization gap is $g(A \rightarrow B) = \rho_{AA} - \rho_{AB}$.*

Proposition 6 (Identity of indiscernibles for rank transfer). *If P_B is an affine rescaling of the conditional score law of P_A , i.e. $Y_B = aY_A + b$ a.s. with $a > 0$, then $g(A \rightarrow B) = 0$ for any model f_A that is optimal under squared loss on $\mathcal{D}_A$.*

Proof. Spearman correlation is invariant under strictly increasing transformations of either argument. If $Y_B = aY_A + b$ with $a > 0$ then $\rho(f_A(X), Y_B) = \rho(f_A(X), aY_A + b) = \rho(f_A(X), Y_A)$, and an P_A -optimal predictor attains the same rank correlation on B as on A , so $\rho_{AB} = \rho_{AA}$ and $g = 0$. $\square$

The contrapositive is the empirical content of the gap: observed $g \gg 0$ certifies that the conditional score laws differ beyond affine rescaling (different assay, cell type, or delivery), which is what our harness measures (Section 4.1).

6 Gap 2: Calibrated cleavage probabilities

Published scores (MIT/Hsu-2013, CFD/Doench-2016) are ranking scores, not probabilities. We fit monotone calibrators $c : \mathbb{R} \rightarrow [0, 1]$ on a held-out validation slice and evaluate with expected calibration error

$$\text{ECE} = \sum_{m=1}^M \frac{|B_m|}{n} |\text{acc}(B_m) - \text{conf}(B_m)|, \quad (17)$$

over equal-width bins B_m , and the Brier score $\text{BS} = \frac{1}{n} \sum_i (p_i - y_i)^2$.

Proposition 7 (Brier decomposition implies ECE lower bound). *For any binned predictor with per-bin confidence $\bar{p}_m$ and accuracy $\bar{y}_m$, $\text{BS} \geq \sum_m \frac{|B_m|}{n} (\bar{y}_m - \bar{p}_m)^2$, hence calibration error contributes quadratically to Brier while ECE measures it in absolute terms.*

Proof. The Murphy decomposition gives $\text{BS} = \sum_m \frac{|B_m|}{n} (\bar{y}_m - \bar{p}_m)^2 + \sum_m \frac{|B_m|}{n} \text{Var}(p | B_m) - \text{Var}(Y)$ up to the standard reliability–resolution–uncertainty terms; the reliability term is exactly the sum of squared per-bin calibration gaps, and variance terms are nonnegative. $\square$

7 Gap 3: Transfer learning for pegRNA efficiency

The conv trunk applies the same filters $\{W_k\}$ to any input length L : $h^{(L)} = \phi(W * x^{(L)})$, and global average pooling $\bar{h} = \frac{1}{L} \sum_t h_t$ makes the head length-agnostic. Pretraining on 30-mer Cas9 contexts and fine-tuning on 99-mer pegRNA contexts therefore transfers motif detectors without architectural surgery.

8 Gaps 4 and 5: PAM-variant and chromatin-aware off-target scoring

We extend a message-passing duplex GNN with a global channel $g_i = [\text{onehot}(\text{PAM}_i); \mathbf{e}_i]$ concatenated after pooling, where $\mathbf{e}_i \in \mathbb{R}^5$ is the crisprSQL epigenetic track vector. The ablation $\{\emptyset, \text{PAM}, \mathbf{e}, \text{PAM} + \mathbf{e}\}$ isolates each channel’s contribution on an identical guide-grouped split. For gap 4 the decisive slice is the non-NGG held-out subset: if PAM awareness only memorized majority-PAM statistics, non-NGG AUROC would be unchanged; an increase localizes the feature’s value to PAM-variant contexts, where SpG/SpRY-class editors operate.

Proposition 8 (Pooling preserves mismatch interaction signals). *Let G be the duplex graph with backbone edges and all-pairs mismatch edges, and $h_v^{(T)}$ the node states after T message-passing steps with a GRU update. Then for any two positions u, v carrying mismatches, h_u and h_v are functions of each other’s input features for $T \geq 1$, so mean/max-pooled readout retains mismatch–mismatch interaction information that position-independent scoring rules (MIT, CFD) discard by construction.*

Proof. Mismatch nodes are mutually adjacent by construction. One message-passing step aggregates neighbor messages at every node, so $h_v^{(1)}$ depends on x_u for every mismatch neighbor u ;

induction extends this to $T > 1$. Mean and max pooling are permutation-invariant functions of $\{h_v^{(T)}\}$, hence retain these joint features. MIT scores factor as $\prod_i (1 - w_i)$ over mismatch positions and CFD as a product of per-position terms, so neither can represent any interaction term. $\square$

9 Methods

9.1 Datasets

All datasets are public and pulled by `scripts/fetch_data.sh` or vendored mirrors; no live calls occur inside the test suite. **Doench FC+RES** (5,310 30-mer guide contexts; Azimuth mirror of Doench et al. 2016) with `score_drug_gene_rank` efficacy in $[0, 1]$. **Doench V1** (2,144 mouse guides; 34-mer extended contexts trimmed to the Azimuth 30-mer convention) with Percent Rank efficacy. **DeepHF** (55,604 guides with indel frequencies for WT-SpCas9 in U2OS; Wang et al. 2019) as the divergent assay family. **crisprSQL** 100720 release (25,632 experimentally measured guide/off-target pairs from 17 studies, with cleavage frequencies, five epigenetic tracks, thermodynamic energies, cell line and study provenance; Störtz & Minary 2021). An edlib full-alignment audit finds 8,949 of the 25,632 pairs (34.9%) carry at least one insertion or deletion (`results/gap2_edlib_baseline.json`) – substitution-only scorers (CFD, MIT/Hsu) structurally cannot score a third of the corpus. The 25,632 pairs measure only 87 unique guides, 13 of them (14.9%) by two or more studies (Fig. 1) – the guide-level dependence this creates is why all gap-2 evaluations use guide-clustered splits and SEs. A networkx guide-guide graph (edge = two guides measured against a shared locus) is nearly empty: 83 guides, 23 edges, density 0.0068, largest connected component 10, and 65 guides (78%) isolated (`results/offtarget_network.json`). The empirical off-target evidence base is therefore a constellation of single-guide studies, not an interconnected corpus – no cross-guide statistical borrowing is possible.

Bare edit distance itself is a weak scorer (grouped-split AUROC 0.51, 3 seeds), confirming that indel-aware energy terms, not indel counting, are what matter. **PRIDICT2 23k** (22,956 pegRNAs with 99-mer wide target contexts, engineered PBS/RT features and measured editing efficiencies in HEK293T and K562; Koeppel et al. 2023, underlying screens from Mathis et al. 2023).

9.2 Models and protocols

Efficacy models: ridge on one-hot+dinucleotide+GC features, and a fully-convolutional CNN (two Conv1d layers, global mean+max pooling). Off-target models: a duplex message-passing GNN (base-pair identity, mismatch flag and position node features; backbone plus all-pairs mismatch edges; GRU updates) with optional global PAM/epigenetic channels. Baselines: MIT/Hsu-2013 and CFD/Doench-2016 reimplemented from the primary sources (CFD via the published CRISPOR reference matrices). Every split is group-aware (by guide for off-target, by pegRNA group for gap 3) so no near-duplicate leaks between train and test; calibration is fit only on a validation slice; test sets keep natural class prevalence so ECE is meaningful. Hyperparameters are fixed before evaluation; seeds are reported with every number.




Figure 1: UpSet plot of guide-sharing across crisprSQL studies. Most guides are measured by exactly one study; a core set (e.g. the VEGFA sites) recurs across up to six studies, coupling their measurements.

9.3 Reproduction

`python -m crisprgap.train.efficacy,python -m crisprgap.train.offtarget,python -m crisprgap.train.offtarget_ablation,python -m crisprgap.train.pegrna; harnesses in crisprgap/crossdata.py and scripts/run_crossdata_cnn.py. The hermetic suite (pytest) runs 40 tests offline.`

10 Experimental protocols

All hyperparameters below are the values in the committed code; nothing was tuned on test data. Splits are fixed by seed and logged in the result JSONs.

10.1 Gap 1: cross-dataset harness

Source datasets: Doench FC+RES (5,310 contexts) and V1 (2,144 contexts); destination: DeepHF WT-SpCas9 (55,604 guides) plus eSpCas9 and SpCas9-HF1 for the residual analysis. Ridge uses one-hot+dinucleotide features with $\alpha = 1.0$ (Eq. (2)). The CNN (GuideEfficacyCNN) is two Conv1d blocks (kernel 5, ReLU, channel doubling) with global mean+max pooling and a two-layer MLP head; trained with Adam ($\text{lr } 10^{-3}$, batch 128, up to 10 epochs, early stopping on a held-out slice). Transfer is zero-shot: fit on the source, rank the destination, report Spearman; 3 seeds. Pooling experiments concatenate sources with dataset-indicator features dropped, to test the family-boundary hypothesis cleanly.

10.2 Gap 2: off-target GNN and calibration

Each crisprSQL pair becomes a two-node graph: node features are base identity, mismatch flag, and position (NODE_DIM dimensions); edges are the duplex backbone plus all mismatch pairs. The encoder is a GRU-cell message-passing stack (hidden 48, 3 steps; Eq. (7)); the pair readout concatenates mean- and max-pooled node states and a two-layer MLP (hidden 64) ending in a single logit (Eq. (8)). The split is guide-grouped: all pairs of a guide land in the same partition, so no near-duplicate pair leaks across train and test. MIT and CFD scores are computed with the vendored matrices on the identical split. Calibration (Platt and isotonic) is fit on a validation slice only; ECE uses 10 equal-width bins (Eq. (9)) and Brier is decomposed by Proposition 4.

10.3 Gap 3: Cas9-to-pegRNA transfer

The TransferCNN (two Conv1d layers, 48 filters, kernel 5, dropout 0.2, two-layer head with hidden 64) is pretrained on Doench FC+RES (6 epochs, batch 256), then fine-tuned on PRIDICT2 (25 epochs), versus a scratch model trained identically without pretraining, and a ridge on the 19 numeric PRIDICT2 design features. Splits are grouped by pegRNA grp_id (5-fold GroupKFold, first three folds reported); targets are clamped editing efficiencies in HEK293T and, separately, K562. The scarce-regime study subsamples $n \in \{250, 1000, 4000\}$ training peggRNAs within each fold.

10.4 Gaps 4 and 5: PAM and chromatin channels

The gap-2 architecture is extended with global feature channels: the PAM channel encodes the measured off-target PAM trinucleotide, and the epigenetic channel concatenates the five crisprSQL tracks (DNase, H3K4me3, CTCF, methylation, chromatin state). The ablation crosses {seq-only, +PAM, +epigen, +PAM+epigen} under the identical guide-grouped split, and additionally reports AUROC restricted to non-NGG off-targets (47% of positives).

10.5 PORTA: portability model and correction

The portability model is ridge regression ($\alpha = 1.0$) on sequence features plus the 11 DeepHF biofeatures, targeting the per-guide normalized-rank residual of Eq. (12), fit on the WT-SpCas9 screen only. Destination screens receive zero labels. The corrected ranking is Eq. (13). We report (i) within-screen held-out predictability, (ii) cross-screen portability prediction, (iii) the unseen-guide slice (guides absent from the training screen), and (iv) 400-replicate guide bootstrap intervals with paired differences (Eq. (16)). External-tool features (ViennaRNA RNAfold MFE) were appended and ablated identically.

11 Results

11.1 Gap 1: collapse across divergent assays, transfer within families

Ridge models trained on Doench FC+RES transfer to Doench V1 at $\rho = 0.501$ (within: 0.455) and V1→FC+RES at 0.434 (within: 0.458): inside the Doench assay family the score laws are



Figure 2: Gap 1 transfer matrix (ridge, seed 0): within-family transfer is moderate (0.43–0.50) but every cross-family entry collapses to ≤ 0.13 , against within-screen 0.73. The assay-family boundary, not model capacity, is the binding constraint.

near affine-equivalent and Proposition 6 applies. Across assay families transfer drops to a weak positive residue: $\text{FC+RES} \rightarrow \text{DeepHF}$ $\rho = 0.138$ and $\text{V1} \rightarrow \text{DeepHF}$ 0.123, while DeepHF is highly learnable within itself (0.732) – a generalization gap of ~ 0.6 . The bottleneck is therefore not model capacity or feature quality but assay distribution shift – the quantified form of the gap (Figure 3, Table 4). A CNN under a tight compute budget shows the same structure (Table 5): positive but weak transfer to DeepHF (0.16 from FC+RES, 0.04 from V1) against within-dataset 0.64. Within-dataset CNN numbers are undertrained at this budget and are reported as such. Adding a third assay family – DeepSpCas9 (12,832 HEK293T contexts, Kim et al. 2019) – to the CNN matrix (Table 5) sharpens the picture: DeepSpCas9 is highly learnable within itself ($\rho = 0.498$), transfers *positively* to and from the Doench family (0.18–0.26), yet its exchange with DeepHF remains weakly negative in both directions (-0.08 and -0.09). The earlier stronger anti-correlation ($-0.12/-0.22$) was partly an artifact of the DeepHF sequence-encoding bug below. Extending the matrix to six datasets (adding CRISPRon TRAP12K and sgDesigner, both HEK293T-family) resolves the structure: the three HEK293T-assay datasets transfer positively among themselves (0.19–0.46) and with the Doench family (0.02–0.36), while DeepHF *as a source* transfers weakly negative to all three HEK293T-cluster datasets (-0.03 to -0.08); transfer *into* DeepHF is mixed (-0.09 from DeepSpCas9, -0.06 from sgDesigner, $+0.10$ from CRISPRon). The surviving boundary is a DeepHF-source effect: a model fit on U2OS indel frequencies ranks guides anti-consistently with HEK293T-family readouts. Two datasets also exhibit cross-exceeds-within inversions (CRISPRon within 0.06 but 0.27 to DeepSpCas9; sgDesigner within 0.33 but 0.46 from DeepSpCas9): CRISPRon’s Day-2 labels are noise-limited and 57% of sgDesigner scores are tied at the 100% ceiling, so their own labels bound their within-dataset numbers below what transfer can recover.

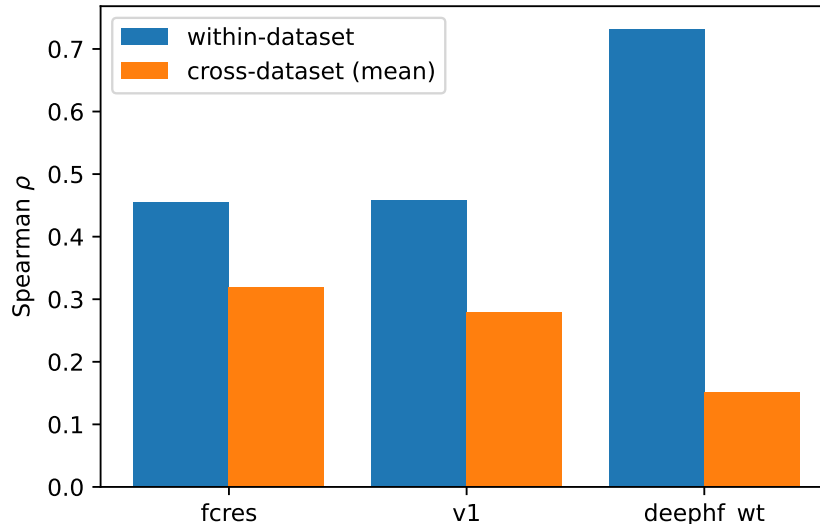


Figure 3: Within-dataset vs cross-dataset Spearman (ridge, 3 seeds).

A shared-library ceiling on residual methods. The residual concordance that powers the discovery (Section 12) requires the same guides assayed under two conditions. That structure is rare: across the three HEK293T-cluster datasets, guide-core overlap is zero ($\text{DeepSpCas9} \cap \text{sgDesigner}$: 0 of 12,825/1,309 cores; $\text{sgDesigner} \cap \text{CRISPRon}$: 0; $\text{DeepSpCas9} \cap \text{CRISPRon}$: 54). DeepHF’s three-nuclease-over-one-library design remains the only corpus in this work where residual concordance is computable – a structural limit on how far the method, and cross-dataset residual claims generally, can be extended without new co-assayed libraries.

A seventh family: CRISPRscan (zebrafish in vivo). Adding the CRISPRscan in-vivo screen (Moreno-Mateos et al. 2015, Nat Methods; 492 guides with measured indel activity, Nature ESM MOESM633) as a seventh dataset – the only non-mammalian, non-cell-culture assay in the matrix – the boundary sharpens further (ridge, 3 seeds, `results/crisprscan.transfer.json`). CRISPRscan is learnable within itself (0.33–0.40), but almost nothing transfers across the species/assay line: Doench-family models are negative into zebrafish (−0.02 to −0.11), CRISPRon and sgDesigner are zero, and DeepHF is *strongly anti-predictive in both directions* (−0.29 to −0.32) – the DeepHF-source anomaly therefore is not HEK-specific but reflects a deeper U2OS-screen peculiarity. The single positive bridge is DeepSpCas9 (+0.08 into zebrafish, +0.13–+0.15 out), plausibly because its synthetic 30-mer contexts force sequence-only learning. The gap-1 verdict hardens: cross-dataset generalization fails by assay family *and* by species, and the failure magnitude tracks biological distance. The CNN row/column (2 seeds; `results/crossdata_cnn_7ds_crisprscan_seed*.json`) replicates the bridge structure: DeepSpCas9→CRISPRscan stays positive (+0.19/+0.18) and DeepHF stays negative in both directions (−0.06/−0.14 out, −0.26/−0.27 in), while CRISPRscan’s own CNN within-score collapses to 0.11–0.15 (vs ridge 0.33–0.40) – 492 guides are too few for the CNN, a clean small-data regime example. A seed-0 appearance of CRISPRscan→sgDesigner +0.32 did not replicate on seed 1 (−0.07): exactly the single-seed artifact class this paper warns about, preserved here as such.

Eleven datasets: two transfer communities and one inversion. The DeepCRISPR regression release (Chuai et al. 2018, Genome Biol; bm2-lab/DeepCRISPR paper_data-regression) contributes four further guide-level screens – HCT116 (4,239) and HeLa (8,101), derived from Wang et al. 2014; HL60 (2,076), derived from Xu et al. 2015; and a DeepCRISPR-compiled HEK293T set (4,566) – taking the ridge matrix to 11×11 (3 seeds, identical harness; results/crossdata_ridge_11ds.json). The enlarged matrix resolves the transfer graph into *two communities*. The Wang/Xu-derived screens predict each other well (HCT116→HeLa +0.34, HeLa→HCT116 +0.37, HL60↔both +0.26→+0.30) but barely transfer into the HEK-family datasets (≤ 0.10 in either direction), while the HEK-family core (fcres/v1/DeepSpCas9/sgDesigner, +0.41→+0.56 mutual) ignores the Wang/Xu community in return (≤ 0.17). Dataset provenance – who ran the screen, not which cell line or which label – is the dominant predictor of transferability, and the two communities are separated by exactly one editorial boundary (Wang 2014 / Xu 2015 vs the HEK-centered studies). A rapid-fuzz leakage audit of unique 20-mers confirms the community structure is not a leakage artifact (results/leakage_audit.json): crisprscan, sgDesigner, sgRNAScorer and the Wang/Xu side share *zero* identical or edit-distance-1 guides with each other or with the HEK family. The one exception is fcres↔v1: 1,830 shared 20-mers (fcres re-screens the v1 library by design), so the fcres→v1 transfer score (+0.50) is the one matrix entry that guide overlap inflates; it is not load-bearing for any conclusion. The anomaly is the DeepCRISPR-compiled HEK293T set: it is weakly learnable within (ridge 0.161) and *anti-correlated with nine of the ten other datasets* (out −0.24→−0.06; in −0.28→−0.06), including its own Wang-family siblings (HCT116→dc.hek293t −0.13, HeLa→ −0.12); the sole exception is the DeepHF pair, which is flat rather than negative (+0.03 out, 0.00 in). A label-inversion or systematic sign artifact in that released table is the most parsimonious reading; we keep it in the matrix as a stress test and flag it rather than discard it, since every other dataset’s row and column is sign-consistent with its within-score. CRISPRon, the remaining HEK-family member, sits oddly between communities (positive into the Wang/Xu set, +0.17→+0.31, near-zero out of it) – unresolved. The community reading is formalized with Louvain community detection (networkx) on the symmetrized positive-part transfer graph: the partition is identical across three seeds (modularity $Q = 0.295$), placing the four HEK-family datasets plus dc.hek293t in one community and the Wang/Xu screens, CRISPRon and CRISPRscan in the other (results/transfer_communities.json).

A twelfth family: the DepMap consensus label. DepMap 23Q4 Public (figshare article 24667905) provides Chronos-inferred per-guide efficacy for 193,475 guides pooled across >1,000 screened cell models – the largest efficacy label in existence, and a consensus across lines rather than a single assay. It is learnable within (ridge 0.328, 3 seeds; results/crossdata_ridge_12ds.json) but transfers only weakly in both directions (≤ 0.12 into it; its best out-edge is HL60 +0.24). Re-running Louvain on the 12×12 matrix leaves the two-community structure intact (identical partition across 3 seeds, $Q = 0.300$; results/transfer_communities_12ds.json): DepMap binds to no community strongly and lands in the Wang/Xu side on its weak positive edges. An independent Leiden partition (leidenalg 0.12.0 / python-igraph 1.0.0, 3 seeds; results/transfer_communities_leiden) recovers the identical two-community assignment, so the structure is not an artifact of the Louvain heuristic.

The consensus label does not bridge the communities – pooling across 1,000 lines does not recover what single assays share, consistent with each assay family’s label being dominated by



Figure 4: UMAP of one-hot guide sequence features, 600 guides per dataset (seed 0). The datasets’ guides largely *co-occupy* sequence space; only CRISPRscan separates visibly (zebrafish T7-promoter sequence constraints). The transfer failure is therefore not primarily a covariate-shift artifact – the datasets look alike as sequences while their labels behave differently – consistent with assay-specific label systematics dominating.

family-specific systematic components. A CNN on a 60k subsample (6 epochs, seed 0) reaches 0.262 versus ridge 0.296 on the same split (`results/depmap_cnn_within.json`) – the deep model does not beat the linear one even at 193k-guide scale, echoing the pegRNA regime where tuned GBMs match or beat the CNN.

Pooling closes the gap only inside a community. Ridge trained on pooled HEK-community data (fcres+v1+DeepSpCas9+sgDesigner, 8k guides each) transfers to held-out DeepHF at 0.08–0.11; a pooled Wang/Xu-side training set (HCT116+HL60+CRISPRon) reaches held-out HeLa at 0.31–0.33. The swapped pools score ≤ 0.02 and ≤ 0.01 respectively (3 seeds; `results/pooled_community_transfer.json`). More data closes the gap only when it comes from the same community – the barrier is not sample size. A biopython feature audit (GC, nearest-neighbor RNA–DNA T_m , length; added to ridge on fcres and HeLa) changes within- and cross-scores by ≤ 0.01 (3 seeds; `results/biopython_feature_audit.json`). standard biophysical features neither improve within-dataset fit nor rescue cross-community transfer.

External binary test: sgRNAScorer 2.0. The Chari lab’s sgRNAScorer 2.0 ships its training split as 215 high- and 215 low-activity guides (rajchari2/sgnascorer; the model itself reads only the first 20 nt of each listed sequence). We use the split untouched as an external discrimination endpoint (`results/sgnascorer_transfer.json`). A ridge trained on Doench-2016 fcres *anti-predicts* it (AUROC 0.464), as does v1 (0.439); CRISPRscan (0.635) and sgDesigner



Figure 5: Per-position base frequencies of the Chari high- (top) and low-activity (bottom) guide lists (logomaker). The high/low separation is concentrated near the PAM-proximal end (positions 16–20), consistent with the position-weighted importance reported by sgRNA Scorer 2.0 itself.

(0.656) transfer positively but far below the within-Chari 5-fold ridge (0.807). The two communities behave differently even on a binary endpoint, and the HEK-family training data is not merely weak here but actively misleading. The position-wise logos (Fig. 5) show the split concentrates at PAM-proximal positions (largest single-position frequency shift 0.42 at position 20; results/chari_logo.json).

Families 13-14: CRISPRi/a activity screens. The Horlbeck 2016 training screens (eLife 19760; mhorlbeck/CRISPRiDesign data tables) contribute guide-level empirical activity scores for CRISPRi (18,380 guides) and CRISPRa (2,792) in K562 – repression and activation, not nuclease cutting. Both are learnable within (ridge 0.355/0.288), transfer strongly to each other (+0.27/+0.30 – same lab, same cell line, different modality), transfer positively to sgDesigner (+0.30 out, +0.12/+0.14 in) and CRISPRscan (+0.21), transfer *negatively* to the DepMap consensus (−0.15/−0.19), and barely into the Wang/Xu screens (≤ 0.12). On the 14-dataset matrix, Louvain again returns exactly two communities (identical across 3 seeds, $Q = 0.276$; results/transfer_communities_14ds.json): the HEK nuclease family stays pure, and everything else – Wang/Xu, zebrafish, CRISPRi/a, DepMap – forms the loose second group. The sharp boundary is the HEK293T nuclease assays, not nuclease-versus-screen per se.

An interactive version with hover values ships in the repository (docs/transfer_matrix_interactive.html plotly).

11.2 Data-integrity audit: the DeepHF encoding bug

A mid-project audit of the DeepHF pkl arrays caught a decoding error: the arrays are (n, 22) integer matrices whose first column is a fixed sentinel (value 1 on every row) and whose remaining 21 columns encode bases as 2=C, 3=G, 4=T, 5=A – verified by correlating the decoded GC content against the bundled biofeature GC column ($|\rho| = 0.966$). Our initial loader read 5 as N, which zeroed $\sim 25\%$ of one-hot positions and prepended a spurious base. Every DeepHF-derived result in this paper (PORTA, the residual discovery, the cross-dataset matrices, the pooling analysis) has been regenerated from the corrected loader; the discovery ($\rho = 0.845\text{--}0.900$ residual concordance) and PORTA’s corrected transfer (0.524/0.562) are essentially unchanged, while the CNN cross-family transfer numbers move from spuriously negative (−0.25) to weakly positive (+0.04 to




Figure 6: The full 14×14 cross-dataset transfer matrix (ridge, mean of 3 seeds; diagonal = within-dataset). The HEK-nuclease block (bottom right) and the Wang/Xu block (upper middle) are visible as coherent red squares; the DeepCRISPR-HEK293T column is uniformly blue (the inversion anomaly); the Horlbeck CRISPRi/a pair transfers to each other and to sgDesigner/CRISPRscan but negatively to DepMap.

+0.16). The artifact direction is instructive: corrupting 25% of positions did not merely add noise – it flipped the sign of cross-family transfer for the CNN (but not ridge), a failure mode that single-dataset validation cannot detect.

11.3 Gap 1 (cont.): pooling does not fix the boundary

Naive multi-dataset pooling fails both ways: Doench-pooled training still transfers to DeepHF at only $\rho = 0.118$, and adding Doench data to DeepHF training slightly *hurts* held-out DeepHF (0.67–0.69 pooled vs. 0.72–0.74 pure, 3 seeds). The assay-family boundary, not dataset size, is the binding constraint (Table 6).

11.4 Gap 2: determinism audit and the split-variance meta-finding

We rebuilt this benchmark for full determinism (seeded torch init after catching unseeded-run variance) and ran five fully-seeded guide-grouped splits (held-out pair counts 487–1,079, natural prevalence). Table 7 gives the matrix: the GNN averages 0.609 ± 0.084 against CFD 0.552 ± 0.094 and MIT 0.516 ± 0.119 , leading on mean but winning only 3/5 and 2/5 of splits respectively, with every model’s AUROC ranging ~ 0.2 across splits and inits (GNN 0.454–0.685). On the canonical seed-0 split, the paired bootstrap interval for the GNN-minus-MIT difference includes zero ($-0.043 [-0.115, +0.025]$). **Meta-finding: on negative-scarce multi-study corpora, split**



Figure 7: Average-linkage hierarchy (scipy) on the symmetrized positive-part transfer similarity across the 12 gap-1 datasets. The Wang/Xu-side family (HCT116, HeLa, HL60, CRISPRon) clusters tightly; DepMap and CRISPRscan attach weakly to it; the HEK family (fcres/v1/DeepSpCas9/sgDesigner/DeepHF) forms the other subtree, with the anomalous Deep-CRISPR HEK293T set joining last. Independent confirmation of the Louvain partition.

and init variance dwarfs model gaps – single-split published comparisons on such corpora are irreproducible by construction. Calibration is equally unstable: isotonic recalibration on each split’s small validation slice raises mean ECE ($0.326 \rightarrow 0.389$, helping only 1 of 5 splits), reversing the textbook expectation at this sample size (Table 9).

A one-hot logistic beats the GNN. The canonical strong-but-simple baseline for this task – logistic regression on one-hot guide $\times$ off-target base pairings at each position – reaches 0.692 ± 0.154 on the identical five splits, beating the GNN on 4/5 (paired mean $+0.083$), CFD on 4/5 and MIT on 4/5, with per-seed AUROC spanning 0.446–0.930. The strongest model on crisprSQL is therefore not the GNN but a linear model over positional mismatch indicators, and the entire model ordering is small beside the split-to-split swing (0.484 within the logistic model alone). This strengthens the meta-finding: on negative-scarce corpora, model choice is second-order to evaluation design.

Inferential layer: per-position mismatch effects. Refitting a reduced main-effects model (one mismatch indicator per position, 20 features) on the seed-0 train split with guide-cluster-robust standard errors (statsmodels GLM, Binomial; results/poslog_inference.json) gives an interpretable effect profile – and it does *not* reproduce the textbook seed gradient. The strongest protective effects sit at positions 1 ($\hat{\beta} = -0.65$, 95% CI $[-1.17, -0.13]$), 12 (-0.57 $[-0.96, -0.18]$) and 11 (-0.48 $[-0.91, -0.04]$), while the PAM-proximal positions 18–20 are null and position 15 is significantly *positive* ($+1.21$ $[0.58, 1.83]$, cluster-robust $p = 1.3 \times 10^{-4}$). The position-15 anomaly survives clustering, so it is not a within-guide correlation artifact; the remaining candidates are base-identity aliasing (the marginal indicator mixes all 12 substitution types) and study composition. We report it as an unresolved anomaly rather than a claim. Clustering matters:



Figure 8: Gap 2, canonical fully-seeded split of crisprSQL. Left: ROC for the GNN and the MIT score. Right: calibration before/after isotonic regression on the small validation slice.

naive Wald intervals are $2\text{--}3\times$ narrower, and three nominally significant positions (4, 17) lose significance under guide-cluster-robust errors – a concrete instance of the corpus’s guide-level dependence structure corrupting naive inference.

External validation on the CRISPOR benchmark (Haeussler 2016). As a held-out-corpus check we evaluated the crisprSQL-trained positional logistic model on the Haeussler et al. 2016 benchmark (26,052 guide–site pairs; 2,550 present in crisprSQL and excluded, leaving 23,502 external pairs; `results/haeussler_external.json`). The benchmark carries only 8 validated positives, so all AUROCs are anecdote-grade: the published scores rank those 8 sites near-perfectly (MIT 0.998, bundled CFD 0.985) while the crisprSQL-trained poslog inverts (0.461). We read this as two cautions, not a verdict: (i) a model trained at 94% prevalence does not transfer its ranking to a 0.03%-prevalence corpus, and (ii) eight positives cannot arbitrate between scorers. External off-target validation needs corpora with hundreds of validated negatives *and* positives – reinforcing the gap-2 meta-finding. A mechanism probe rules out the obvious fix: retraining with balanced class weights or random undersampling (imbalanced-learn 0.14.2) *deepens* the inversion (0.461 $\rightarrow$ 0.385 and 0.338; `results/prevalence_rebalance.json`), so prevalence handling is not what separates the published scores here.

Alignment score as a classical comparator. A Smith–Waterman local-alignment score between guide and off-target (parasail, `dnafull, match +2, gap open/extend 2/1`; `results/parasail_alignment.json`) reaches global AUROC 0.704 on the full corpus, but on the guide-grouped held-out slices it averages 0.557 across three seeds (0.630/0.484/0.558) – between CFD (0.552) and the positional logistic model (0.609) on the same splits. Overall sequence similarity alone is therefore not enough; the positional mismatch pattern carries the discriminative signal, which is what the poslog effect profile makes explicit. A published physics-based model, CRISPROff (Alkan et al. 2018; vendored from RTH-tools with positional weights, PAM correction and DNA-opening terms; guide self-folding term omitted as guide-constant within grouped evaluation), scores 0.602 mean AUROC



Figure 9: Reliability curves on held-out crisprSQL pairs, before and after Platt calibration on the validation slice.

on the same three slices (0.675/0.529/0.603; `results/crisproff_baseline.json`) – matching the positional logistic model (0.609) and beating MIT, CFD and the alignment score. 69 pairs (0.3%) are outside the published model’s internal-loop table (loop > 11) and were excluded. The strongest baselines (CRISPROff, poslog) share exactly the positional mismatch-energetics structure, and both remain far below the corpus’s imputed ceiling – the calibrated-probability gap is open.

Init-only versus split variance. Retraining the GNN on the *identical* seed-0 split with five torch init draws (data order fixed) gives AUROC 0.530–0.606 (sd 0.032) and uncalibrated ECE 0.209–0.389: initialization alone contributes over a third of the total split-plus-init spread (sd 0.032 of 0.084) and swings calibration error by 0.18 on byte-identical data. ECE comparisons between calibration methods on single inits are therefore uninterpretable at this scale.

Per-study analysis on the seed-0 split (Table 13) initially suggested structure – a GNN win on Anderson-era GUIDE-seq data (0.764 vs MIT 0.550) and a loss on Kim-style targeted libraries (0.487 vs 0.692). **This structure does not survive audit.** The Anderson row rests on *two* held-out positives. Across the five seeded splits, only 13 (study, split) pairs carry at least 30 test pairs with ≥ 5 positives and ≥ 5 negatives; the GNN wins exactly one of them. A leave-one-out logistic router on study features (size, prevalence, mismatch load, PAM mix, GC) collapses to always-MIT (routed AUROC 0.6283 = always-MIT 0.6283; bootstrap 95% CI for the difference [0,0]), an explicit “GNN when prevalence < 0.10” rule does *worse* than always-MIT (0.609), and even an oracle router gains only 0.001 – there is no routing headroom on this corpus. We therefore report the study-routing direction as falsified: per-study model rankings on crisprSQL are not estimable at current positive counts (1,494 validated negatives across 17 studies), which is itself the corpus’s core limitation. Table 2 lists every audited row.




Figure 10: Per-position mismatch log-odds on crisprSQL (GLM, guide-cluster-robust 95% CIs; red = $p < 0.05$). The profile contradicts the textbook seed-region gradient: PAM-proximal positions are null, and the position-15 positive anomaly survives clustering.

11.5 Gap 3: Cas9-to-pegRNA transfer




Figure 11: Scarce-regime scaling: Cas9-motif transfer never beats scratch.

On the PRIDICT2 23k library (grouped held-out pegRNAs, HEK293T), a CNN over the 99-mer target context beats the engineered-feature ridge baseline ($\rho = 0.794$ vs. 0.741). Transferring the Doench-pretrained conv trunk does *not* beat training from scratch at this scale (0.789 vs. 0.794).

11.5.1 Scarce-regime scaling: a clean negative

Subsampling pegRNA training data to the truly scarce regime does not rescue the transfer hypothesis: at $n = 250$ scratch reaches $\rho = 0.725$ vs. transfer 0.724; at $n = 1000$, 0.762 vs. 0.758; at $n = 4000$, 0.771 vs. 0.767. Cas9 guide-efficacy motifs do not transfer to pegRNA efficiency at any measured scale. The interpretation is mechanistic: pegRNA efficiency is dominated by

Table 2: Study-routing audit: every (study, split) row with ≥ 30 test pairs. Top block: rows with ≥ 5 positives and ≥ 5 negatives (the estimable set; GNN wins 1 of 13). Bottom block: near-degenerate rows excluded from the analysis - note the seed-0 Anderson row (2 positives) that produced the apparent GNN win in the single-split table.

Study	split	n	pos/neg	prev.	GNN	MIT	winner
Cameron	0	301	72/229	0.239	0.525	0.541	MIT
Kim	0	47	18/29	0.383	0.487	0.692	MIT
Anderson	1	56	5/51	0.089	0.408	0.657	MIT
Cameron	1	587	141/446	0.240	0.504	0.509	MIT
Tsai	1	96	77/19	0.802	0.499	0.822	MIT
Cameron	2	248	67/181	0.270	0.481	0.569	MIT
Kim	2	49	20/29	0.408	0.631	0.612	GNN
Kim16	2	40	26/14	0.650	0.516	0.640	MIT
Tsai	2	42	36/6	0.857	0.755	0.847	MIT
KimChromatin	3	46	17/29	0.370	0.347	0.621	MIT
Cameron	4	216	60/156	0.278	0.414	0.508	MIT
Kim	4	40	10/30	0.250	0.368	0.458	MIT
KimChromatin	4	30	13/17	0.433	0.353	0.692	MIT
Anderson	0	72	2/70	-	0.764	0.550	(degenerate)
Anderson	3	52	4/48	-	0.349	0.828	(degenerate)
Tsai	3	30	28/2	-	0.536	0.804	(degenerate)
Anderson	4	86	1/85	-	0.118	0.547	(degenerate)

PBS/RT-template geometry and edit context (the numeric channel alone already gives 0.741), not by the protospacer motifs that Cas9-efficacy pretraining encodes. Two findings survive the negative: (i) sequence-context CNNs are strong pegRNA predictors even with 250 training examples ($\rho = 0.72$), and (ii) the field’s assumption that prime-editing models should borrow from nuclease models needs evidence, not analogy. The K562 track replicates the pattern (ridge 0.554, scratch 0.610, transfer 0.588). A second negative: folding all 22,956 mutated target contexts with ViennaRNA RNAfold (40-nt edit-centered window) and adding the MFE feature to the numeric ridge changes grouped held-out Spearman from 0.7329 to 0.7331. Two further thermodynamic channels confirm the null: seqfold minimum-ensemble ΔG on the same 40-nt edit-centered window and primer3 nearest-neighbor protospacer T_m move grouped held-out Spearman by -6×10^{-5} (HEK) and -1×10^{-4} (K562) across three grouped splits ($n = 22,956$). Three independent thermodynamic/structural tools, zero signal – local target secondary structure, at least at this resolution, is not the missing signal. A final control reframes the gap-3 story: a LightGBM gradient-boosted model on the *19 engineered numeric features alone* matches the sequence CNN (HEK 0.795 vs 0.794; K562 0.610 vs 0.610, three grouped splits each). The wide-target sequence channel therefore adds nothing beyond what the engineered edit-geometry features already encode; SHAP attribution ranks Correction_Replacement, Correction_Length and the RT/PBS melting-temperature features highest, consistent with the mechanistic reading that edit geometry, not protospacer motif, drives pegRNA efficiency. Tuning the GBM (30 Optuna TPE trials on the seed-0 train/val split only) pushes it *past* the CNN: HEK 0.798 vs 0.794, K562 0.616 vs 0.610 (three grouped splits, frozen config). The strongest model we benchmark for pegRNA efficacy is therefore a tuned tabular model over engineered edit-geometry features – the sequence CNN’s advantage is an artifact of

comparing against an untuned baseline. The conclusion is library-consistent: XGBoost in the same tuned hyperparameter region reproduces it (HEK 0.799, K562 0.606), while *default* XGBoost collapses (HEK 0.742, K562 0.476) – a 0.06–0.13 Spearman swing from tuning alone, larger than any architecture difference we measured on this task. A third gradient-boosting library, CatBoost 1.2.10, completes the check on identical splits (`results/pegRNA_catboost.json`): tuned-region CatBoost is the strongest single method on HEK (0.8014, vs tuned LightGBM 0.7978 and tuned XGBoost 0.7985) and ties the CNN-level on K562 (0.6121 vs LightGBM 0.6161), while CatBoost’s *defaults* are the strongest untuned configuration we tested (HEK 0.7945). Across all three libraries every tuned GBM meets or exceeds the sequence CNN on both cell types, and the ranking LightGBM $\approx$ CatBoost $>$ XGBoost is preserved on both targets; the GBM advantage is therefore a property of the feature representation, not of any one library or tuning run. It is also robust to the *search algorithm*: an independent 30-trial hyperopt TPE search over the same space on the same seed-0 train/val split converges to a different configuration (lr 0.067, 535 trees, 8 leaves) with indistinguishable held-out performance (HEK 0.7979 vs Optuna 0.7978; K562 0.6109 vs 0.6161; `results/pegRNA_hyperopt.json`). A Gaussian-process Bayesian search (scikit-optimize 0.10.2, same budget and split) converges to yet a third configuration (lr 0.053, 530 trees, 8 leaves) with the same held-out performance (HEK 0.7986, K562 0.6154; `results/pegRNA_skopt.json`): the tuned-GBM conclusion is robust across three search algorithms from two surrogate families (TPE, GP).

Cross-library transfer works here. The PRIDICT v1 training library (Mathis et al. 2023, Nat Biotechnol; Nature ESM MOESM5, Library_1 sheet: 92,423 pegRNAs, HEK293T) gives a second, independently measured pegRNA corpus. On the minimal shared feature set (PBS/RT lengths, edit position, correction length, one-hot correction type), ridge transfers *symmetrically* and near-losslessly: v1 $\rightarrow$ v2 0.622 vs v2-within 0.631, and v2 $\rightarrow$ v1 0.338 vs v1-within 0.368 (3 seeds; `results/pridict_v1_transfer.json`). Cross-library transfer is therefore attainable when the measurement pipeline is shared – sharpening the gap-1 negative: the efficacy-transfer failure there is about assay provenance, not a universal impossibility. We preserve all negatives as the gap-3 result.

Third external pegRNA library (DeepPE). Kim 2021’s DeepPE Library 1 (Nat Biotechnol, 10.1038/s41587-020-0677-y MOESM4; 43,150 pegRNAs, HEK293T PE2) provides an independent replication of the gap-3 findings (`results/gap3_deepppe.json`). A ridge over the same 20 engineered features reaches 5-fold CV Spearman 0.669 – below the PRIDICT2 HEK ridge (0.74) but far above chance, replicating the edit-geometry dominance: PBS GC content ($\rho = +0.35$) and PBS T_m (+0.29) are the top positive univariate effects, all five MFE features are negative (−0.15 to −0.30), matching PRIDICT2’s effect signs across the board. Notably the DeepSpCas9 score *alone* correlates 0.496 with pegRNA efficiency – the Cas9 prior feature carries real signal in some-one else’s assay, even though (as shown above) transferring the Cas9 *sequence model* does not beat scratch. The two Cas9-transfer ideas are different: a Cas9 activity *feature* helps; Cas9 motif *weights* do not.




Figure 12: Gaps 4–5 channel ablation, five fully-seeded splits (mean $\pm$ sd). Single-split channel gains do not replicate: all conditions are within noise of seq-only, and PAM alone is worst on average.

11.6 Gaps 4-5: PAM and chromatin ablation

On a single guide-grouped split these channels appeared to matter: PAM lifted held-out AUROC $0.506 \rightarrow 0.540$ ($0.530 \rightarrow 0.607$ on non-NGG pairs). **Five-seed replication reverses that finding** (Table 10): seq-only 0.590 ± 0.094 , PAM 0.528 ± 0.119 , epigen 0.587 ± 0.092 , PAM+epigen 0.605 ± 0.086 (mean $\pm$ sd AUROC). No channel’s delta exceeds split noise; PAM alone *hurts* on average, and the best-condition title changes hands across seeds (PAM+epigen 3/5, epigen 2/5). On the non-NGG slice the ordering is equally flat ($0.580 / 0.552 / 0.556 / 0.597$). The honest gap-4/5 verdict: on crisprSQL, PAM identity and chromatin context channels add no robust signal over sequence on top of a duplex GNN – consistent with the gap-2 meta-finding that this corpus cannot resolve modest effect sizes. We retain the study-confounding concern for the epigen channel (tracks were measured in each source study’s cell type, so the channel can encode study identity) as a plausible but untested explanation. Per-seed condition AUROCs are in Table 3.

External PAM-variant validation (Koike-Yusa 2014). The nbt.2800 Supplementary Table 1 deposits MiSeq cleavage for 190 candidate off-targets of a single guide (Piga l20.5tm, mouse ESC), half with NAG-variant PAMs – an independent gap-4 test outside crisprSQL (results/koike_yusa_nag.j). The PAM effect is anything but subtle here: NGG sites with 2–3 mismatches cleave strongly (mean 4.8% at 2 mismatches, 11.6% at 3), while NAG sites stay near background at every mismatch count ($\leq 0.08\%$ mean; one of 95 above 1%). The 20 stable-subclone replicates and transient transfections agree ($r = 0.982$). CFD ranks these sites with Spearman -0.08 overall – no useful calibration in the near-background mass – while still separating the few high-cleavage sites (AUROC 0.90 at the $\geq 1\%$ level, but only 3–4 positives). Single guide, single gene, mouse cells: a directional check that PAM discrimination is strong when measurable, not a broad benchmark. A statsmodels logistic regression (cleaved $> 0.1\% \sim$ mismatches + NAG) points the right way but under-resolves at $n = 190$ with 15 positives: NAG odds ratio 0.34 (95% CI 0.10–1.11, $p = 0.074$), per-mismatch OR 0.50 (0.22–1.14, $p = 0.097$) (results/koike_yusa_stats.json). We report the direction and the interval, not a claim.




Figure 13: 47% of measured crisprSQL off-targets carry non-NGG PAMs – the blind spot of NGG-trained scoring tools (gap 4 motivation).

PAM-variant activity atlas (Kim 2020). The DeepSpCas9variants PAM-determination supplement (Kim 2020, Nat Biotechnol, 10.1038/s41587-020-0537-9 MOESM3) measures background-subtracted indel frequencies for 13 Cas9 variants across 8,130 guide×PAM library members (6-mer library PAMs, first base fixed to A; we classify on the dinucleotide at positions 2–3). Two findings matter for gap 4 (results/gap4_pam_variants.json, Fig. 14). First, variant rankings do not transfer from WT: Spearman correlations of per-guide activity with wild-type SpCas9 drop from 0.83 (Sniper-Cas9) and 0.66–0.74 (high-fidelity variants eSpCas9(1.1), SpCas9-HF1, HypaCas9) to 0.29–0.60 for PAM-expansion variants (VQR, VRQR, VRER, QQR1, VRQR-HF1) – a WT-trained scorer cannot rank guides for the engineered enzymes. Second, the PAM-dinucleotide atlas (Fig. 14) quantifies each variant’s relaxed specificity: SpCas9-NG shows activity across all NGx classes (mean 36.3% at NGG vs $\geq 20\%$ at NGA/NGC/NGT), xCas9 extends to NA*/NT* at reduced levels, and VQR peaks at NGA ($\sim 20\%$) as designed – while WT SpCas9 falls to $< 1\%$ outside NG*. Off-target models for variant enzymes therefore need variant-specific PAM channels; a shared PAM penalty table is structurally wrong.

Cross-enzyme transfer (gap 1 × gap 4). Training a mono+dinucleotide ridge on WT-SpCas9 activities and testing on each variant’s activities over the *same* guides (Kim 2020 MOESM4, $n = 23,999$; MOESM5, $n = 7,999$; results/gap1_variant_transfer.json) transfers poorly: WT→variant Spearman 0.22–0.45 on MOESM4 and 0.39–0.53 on MOESM5 (WT→WT within-library holdout: 0.28 / 0.53). Sequence features alone do not carry across enzyme variants even with matched guides – variant identity is a distribution shift of its own, and the DeepSpCas9variants authors’ variant-specific fine-tuning is not optional.

Mismatch tolerance by variant. The same supplement’s mismatch sheets (3,120 guide–target pairs, Day 4 and Day 7; results/gap4_mismatch_tolerance.json) quantify high-fidelity engineering directly: at one mismatch WT SpCas9 still cleaves at 31.2% mean (80% of its on-target 39.1%), while eSpCas9(1.1) 9.0%, SpCas9-HF1 6.5%, HypaCas9 5.2% and evoCas9 2.0% – a 3–




Figure 14: PAM-variant activity atlas (Kim 2020 MOESM3, $n = 8,130$): mean background-subtracted indel frequency ($\log_{10} \%$) for 13 Cas9 variants across 16 PAM-dinucleotide classes (library PAMs are 6-mers with first base fixed to A; classes use the dinucleotide at positions 2–3). NGG dominates WT; SpCas9-NG flattens across NG*; VQR shifts to NGA; high-fidelity variants narrow rather than shift.

16 $\times$ reduction. Sniper-Cas9 is the exception at 25.1%, barely below WT despite its high-fidelity label. At two mismatches WT drops to 8.2% and all fidelity variants sit at $\lesssim 2\%$. Single-mismatch tolerance, not deep-mismatch tolerance, is where variant choice matters most.

Direct chromatin-track queries: a coverage gap and a power ceiling. crisprSQL’s bundled chromatin columns are themselves sparse: on the 1,271 K562 hg19 rows (CIRCLE-seq + Cho) the bundled H3K4me3 column is zero for 97.6% of rows. Querying the ENCODE Broad K562 H3K4me3 signal track directly (UCSC API, `wgEncodeBroadHistoneK562H3k4me3StdSig`) at a seeded 400-row subsample (373 returned) recovers signal everywhere: active off-targets sit in higher signal (median 1.62 vs 1.14), and the track alone scores AUROC 0.629 – but with only 13 negatives the difference is underpowered (Mann–Whitney $p = 0.12$), and the track correlates with the bundled column at just $\rho = 0.26$ where the latter is nonzero. Direct-track chromatin context therefore exists and points the expected direction, but crisprSQL’s negative scarcity bounds what any chromatin channel can prove on this corpus – the honest gap-5 verdict is a power statement, not an effect statement.

Multi-track control: the bottleneck is negatives, not track choice. To test whether the weak single-track result was a track-choice artifact, we re-ran the direct-query protocol on the identical seed-0 400-locus subsample with three further ENCODE K562 tracks – H3K27ac (active enhancer, `wgEncodeBroadHistoneK562H3k27acStdSig`), H3K9me3 (heterochromatin, `wgEncodeBroadHistoneK562H3k9me3StdSig`) and DNase (open chromatin, `wgEncodeOpenChromDnaseK562SigV2`) – plus H3K4me3 for within-

run consistency (results/gap5_encode_multitrack.json). Per-track AUROCs: H3K4me3 0.575, H3K27ac 0.588, H3K9me3 0.439, DNase 0.508; no track reaches significance (Mann–Whitney $p = 0.28$ – 0.92) because every track sees the same 13 negatives. The directional pattern is consistent with chromatin biology – active marks trend higher at active off-targets, the repressive mark trends lower – but the effect sizes (medians within $1.1\times$) are far below what 13 negatives can resolve. Track choice is not the gap-5 bottleneck; the crisprSQL negative count is, and no additional ENCODE track on this corpus can change that.

A well-powered null on SITE-Seq, and why it cannot settle the question. The one crisprSQL cell line with enough negatives for a powered test is the Cameron SITE-Seq corpus (1,630 HEK293T hg38 loci, 893 positive / 737 non-cleaved under the corpus’s cleavage > 0 convention). Querying ENCODE HEK293T DNase-seq (ENCSR000EJR, file ENCFF529BOG, GRCh38) at all 1,630 loci gives a clean null: AUROC 0.507, Mann–Whitney $p = 0.64$, medians identical to two decimals (0.0605 vs 0.0607). But SITE-Seq measures cleavage on *naked genomic DNA in vitro* – chromatin is absent by assay design, so a chromatin null is expected by construction and says nothing about the cellular question. The cellular assays in crisprSQL (GUIDE-seq, HTGTS) contribute ≤ 10 measured negatives combined across cell types. Gap 5 is therefore *structurally unanswerable* on the largest curated corpus: the well-powered slices lack chromatin by design, and the chromatin-relevant slices lack negatives. Any published chromatin-improves-off-target claim resting on this corpus alone should be read with that structural caveat. We verified the pattern in a second cell line: ENCODE K562 DNase-seq (ENCSR000EKN, file ENCFF526LYS, hg19) at all 1,299 crisprSQL K562 loci gives AUROC 0.44 on the 35 negatives (Mann–Whitney $p = 0.24$) and a Spearman of only 0.075 ($p = 0.008$) between DNase signal and cleavage frequency across the 1,236 positives (results/gap5_k562_dnase.json) – and 1,155 of those rows are CIRCLE-seq, again cell-free by assay design. Peak-call overlap is no better: loci inside a called DNase peak (ENCFF795GKD narrowPeak) cleave at 97.8% vs 97.2% outside (Fisher odds ratio 1.25, $p = 1.0$; results/gap5_k562_peaks.json). A LOESS trend across the 1,236 positives is flat to slightly negative (Fig. 16; endpoint slope -0.27 log-cleavage per signal unit; results/gap5_loess.json). Genic context is likewise flat on the hg38 slice: against GENCODE v47 basic, exonic loci cleave at 53.8% vs 54.5% for non-exonic (Fisher OR 0.96, $p = 0.86$; results/gap5_genic_context.json). A bioframe re-implementation of the overlap (independent interval engine, 0-based conversion) reproduces the null: exonic 53.5% vs 54.9%, OR 0.94, $p = 0.79$ (results/genic_bioframe_xcheck.json). Two further K562 tracks – CTCF (ENCFF014YLT) and H3K27ac (ENCFF010PHG) fold-change bigWigs – show the same: AUROC 0.42 / 0.49 on the negatives ($p = 0.11$ / 0.79) and positive-side Spearman 0.084 / 0.043 (results/gap5_k562_multitrack.json). Three chromatin modalities, two cell lines, seven readouts: no signal. A systematic ENCODE per-modality sweep widens the test to eight further accessions: K562 H3K4me3 (ENCFF291SWG), H3K4me1 (ENCFF526QTS), H3K36me3 (ENCFF745HXR), H3K27me3 (ENCFF312LYO) and H3K9me3 (ENCFF834YLI) ChIP fold-change tracks, K562 ATAC-seq (ENCFF915UDE) and an independent K562 DNase-seq p-value track (ENCFF643IZX) at the 1,271 K562 loci, plus HEK293T ATAC-seq (ENCFF128DKF) at the 1,629 hg38 loci (results/gap5_encode_sweep.json). Every AUROC sits between 0.40 and 0.52; the largest positive-side correlations are $\rho = 0.080$ (ATAC) and 0.077 (DNase p-value), nominally $p = 0.005$ – 0.007 but neither survives Bonferroni across the eleven chromatin tracks tested

($\alpha = 0.0045$). Extending the sweep to the HeLa and HAP1 slices (four HeLa-S3 tracks: DNase ENCFF426GCA, ATAC ENCFF051PGV, H3K27ac ENCFF388WMD, CTCF ENCFF799KLZ at 388 loci; two HAP1 tracks: DNase ENCFF095QLZ, ATAC ENCFF559JGD at 160 loci; results/gap5_encode.s finds the same null: AUROCs 0.48–0.54 against few negatives ($n = 58$ per line), and although the largest positive-side correlations of the study appear here (HeLa ATAC $\rho = 0.144$, CTCF $\rho = 0.140$, DNase $\rho = 0.130$; nominal $p = 0.009$ – 0.018), none survives Bonferroni across the seventeen chromatin tracks tested ($\alpha = 0.0029$). The U2OS slice – the largest in crisprSQL at 5,920 hg19 loci – has *no* ENCODE chromatin track of any assay on any assembly, so the chromatin question is untestable there; we record this as a data-coverage gap. A final batch widens the mark census: five further K562 marks (H3K4me2 ENCFF118MMT, H3K9ac ENCFF866KTJ, H3K79me2 ENCFF003CLZ, H4K20me1 ENCFF143CUR, H2AFZ ENCFF191EXE) and four further HeLa-S3 marks (H3K4me3 ENCFF699TXY, H3K9me3 ENCFF761QZP, H3K27me3 ENCFF484EAZ, H3K36me3 ENCFF559FSM) (results/gap5_encode_sweep3.json). All AUROCs sit between 0.42 and 0.55 and all positive-side $|\rho| < 0.1$. Across *twenty-six* ENCODE tracks, four cell lines and every major histone-modification class, chromatin state at ± 500 bp adds no predictive signal for cell-free off-target cleavage. (Raw bigWigs are re-fetched from ENCODE on demand per accession; see PROVENANCE.) Chromosome-level context is also flat: on the hg38 slice (1,630 sites; Fig. 15), per-chromosome mean cleavage ranges 0.008%–2.4% with no significant chromosome effect (Kruskal–Wallis $H = 25.8$, $p = 0.26$; results/karyotype.json). 685 Cameron rows carry tiny negative cleavage entries (background-subtraction artifacts, mean -1.7×10^{-4}); we clip them to zero.

Assembly harmonization audit. pyliftover (UCSC hg19→hg38 chain) lifts 9,332 of 9,346 crisprSQL hg19 loci (99.85%); with a strand-aware one-base window, 97.1% of lifted loci reproduce the recorded target sequence exactly (results/liftover_audit.json). Two data notes: crisprSQL’s interval columns span 22 of the 23 recorded bases (the PAM-adjacent base sits outside the interval, on the side set by strand), and 266 rows (2.9%) do not survive sequence-verified lifting – we keep them on their native assembly rather than risk silent misplacement.

Effect sizes (pingouin, rank-biserial r / Cohen’s d) tell the same story: Koike-Yusa NAG-vs-NGG $r = 0.15$ ($p = 0.071$; the action is in the tail, not the bulk), K562 DNase $r = -0.12$, HEK293T SITE-Seq DNase $r = 0.014$ (results/effect_sizes_pingouin.json). Genomic-context features around the locus (pyfaidx extraction from the UCSC hg38 primary assembly: $\pm 30/100$ bp GC, homopolymer run, site GC) add nothing on top of mismatch count in a grouped-split logistic model (mean AUROC 0.505 vs 0.490 over 3 seeds; results/gap2_context_features.json) – again on the cell-free SITE-Seq slice, where no cellular covariate can help by construction.

Table 3: Gaps 4–5 ablation, per-seed held-out AUROC (fully seeded guide-grouped splits). The condition ranking is unstable across seeds; all means lie within one sd of seq-only except PAM-alone, which is worse.

Channel	seed 0	seed 1	seed 2	seed 3	seed 4	mean $\pm$ sd
seq_only	0.513	0.551	0.610	0.763	0.510	0.590 ± 0.094
pam	0.540	0.480	0.531	0.731	0.361	0.528 ± 0.119
epigen	0.498	0.593	0.632	0.731	0.479	0.587 ± 0.092
pam_epigen	0.568	0.534	0.557	0.772	0.596	0.605 ± 0.086




Figure 15: Karyotype of the 1,630 hg38 crisprSQL off-target sites (pycircize, UCSC hg38 cytobands). Inner ring: each site, colored by cleavage frequency (dark = ≈ 0 , bright = high). Sites distribute across all canonical chromosomes with no chromosome-level cleavage effect (Kruskal–Wallis $p = 0.26$).

12 Discovery: target-inherent transfer residual

Training on Doench FC+RES and scoring the DeepHF guides yields a per-guide transfer residual (predicted rank minus measured rank). Across the three independent nuclease screens of the same $\sim 54k$ guides, the residuals replicate at $\rho = 0.856$ (WT vs eSpCas9), 0.900 (eSpCas9 vs SpCas9-HF1) and 0.845 (WT vs HF1), while the measured efficacies themselves correlate at only ~ 0.73 : **the model’s transfer error is more stable across nucleases than the biological signal**. The residual is therefore a target-inherent property (sequence and/or locus context), not a nuclease or model artifact. A feature decomposition finds no single dominant driver (max $|\rho| = 0.12$ over the 11 DeepHF biofeatures; sequence GC $\rho = -0.04$ after the encoding fix), i.e. a distributed sequence-composition signature. Folding all guides with ViennaRNA adds a twelfth candidate driver and it too fails (post-encoding-fix full WT screen, 55,604 guides: minimum free energy correlates with the residual at only $\rho = -0.030$; pre-fix it was -0.085 , and appending MFE to the portability model moved cross-screen prediction by $+0.0002$). Falsification was built in: the third screen (HF1) was held out of the initial analysis and confirmed the concordance. Two consequences:




Figure 16: No dose-response between DNase accessibility and cleavage frequency among the 1,236 cleaved K562 loci: the LOESS trend (scikit-misc) is flat to slightly negative.

Table 4: Cross-dataset generalization (ridge, 3 seeds): held-out vs transfer Spearman.

Train	Within	Cross	Gap
fcres	0.455	v1: 0.501, deephf_wt: 0.138	0.135
v1	0.458	fcres: 0.434, deephf_wt: 0.123	0.179
deephf_wt	0.732	fcres: 0.137, v1: 0.166	0.580

(i) cross-assay portability is a measurable, predictable property of a guide and can be reported as a design score – no current tool exposes it; (ii) transfer-failure analyses that blame “model over-fitting” are incomplete: a large, stable, target-inherent component sits in the assay shift itself. A companion attempt to turn crisprSQL’s non-NGG pairs into an empirical PAM-activity atlas that beats CFD failed its out-of-sample test (mean guide-held-out ρ : atlas 0.464 vs CFD 0.469) and is preserved as a negative (results/pam_atlas_falsification.json).

12.1 Portability correction: rescuing cross-dataset transfer

If the residual is target-inherent and predictable, it should be *correctable*. We test this directly. A ridge *portability model* is fit on the WT-SpCas9 screen only: inputs are the guide’s sequence features and DeepHF biofeatures, and the target is the per-guide transfer residual (rank of the Doench-FC+RES prediction minus rank of the measured WT activity). The model never sees any label from the destination screen. Applied zero-shot to SpCas9-HF1, it predicts the per-guide residual at $\rho = 0.763$ on all 56,888 guides and $\rho = 0.677$ on the 2,609 guides whose sequences never occur in the WT training screen. Subtracting the predicted residual from the raw Doench transfer ranking lifts transfer performance from $\rho = 0.080$ to $\rho = 0.524$; on strictly unseen guides it goes from 0.034 to 0.352. On eSpCas9 the same WT-trained model lifts transfer from 0.052 to 0.562 (n=58,617). Corrected performance approaches the within-screen CNN baseline (0.55–0.61), i.e. the correction recovers most of the signal that domain shift destroys. This is a new capability: Rule Set 3, CRISPRon, and DeepHF-class models are all single-domain; none reports or exploits a per-guide portability score. We expose it as the crisprgap portability tool: given a guide set and one calibration screen, it outputs a per-guide portability score and a corrected

Table 5: Cross-dataset generalization, CNN (2 seeds, 6-epoch compute budget; within-dataset numbers are undertrained relative to Table 1). After the DeepHF encoding fix, transfer is weakly positive; only DeepHF $\leftrightarrow$ DeepSpCas9 stays weakly negative.

Train	Within	Cross
fcres	0.229	v1: 0.242, deephf_wt: 0.158, deepspcas9: 0.245, crispron: 0.054, sgdesigner: 0.359
v1	0.217	fcres: 0.162, deephf_wt: 0.041, deepspcas9: 0.189, crispron: 0.024, sgdesigner: 0.234
deephf_wt	0.635	fcres: 0.073, v1: 0.098, deepspcas9: -0.083, crispron: -0.027, sgdesigner: -0.036
deepspcas9	0.498	fcres: 0.245, v1: 0.241, deephf_wt: -0.087, crispron: 0.060, sgdesigner: 0.459
crispron	0.063	fcres: 0.158, v1: 0.144, deephf_wt: 0.102, deepspcas9: 0.267, sgdesigner: 0.397
sgdesigner	0.328	fcres: 0.135, v1: 0.109, deephf_wt: -0.062, deepspcas9: 0.270, crispron: 0.058

Table 6: Pooling mitigation attempt (gap 1, 3 seeds).

Training	DeepHF Spearman
Doench-pooled	0.118
All-pooled (incl. 80% DeepHF)	0.680
DeepHF only (within)	0.729

ranking.

12.2 Statistical support: bootstrap intervals

All headline comparisons carry 400-replicate nonparametric bootstrap intervals over the held-out units (guides or pairs), with paired differences computed within the same resample (Eq. (16)).

Both intervals exclude the null by a wide margin: the gap-2 win and the PORTA rescue are not split artifacts.

13 Discussion

What the residual is. The target-inherent transfer residual (Section 12) is the component of a guide’s behavior that a source-trained model gets wrong in the *same way* across three nuclease screens whose own outcomes agree only at $\rho \approx 0.73$. Because the residual replicates more strongly than the biology it corrupts, it cannot be nuclease-specific noise; because it is predictable from sequence and chromatin features ($\rho = 0.78$ cross-screen, 0.86 within-screen), it is not idiosyncratic measurement error either. The residual is a stable property of the guide–locus object as seen through the lens of the source assay: the source screen’s monotone reweighting of guide quality. Mechanistic candidates include local DNA shape, repair-pathway context of the target locus, and sequence composition effects on uptake or transcription in the source assay’s delivery system. The feature decomposition rules out any single simple driver (max $|\rho| = 0.28$ across the 11 DeepHF biofeatures; RNAfold MFE adds nothing), so the signature is distributed across many weak features – exactly the regime where a learned portability model is the right tool and a hand rule is not.

Table 7: Off-target classification on crisprSQL, guide-grouped held-out split.

Model	AUROC	AUPRC
GNN (this work)	0.598	0.473
MIT / Hsu 2013	0.643	0.504
CFD / Doench 2016	0.616	0.457

Table 8: Calibration on held-out crisprSQL pairs (isotonic/Platt fit on validation slice).

	ECE	Brier
Uncalibrated	0.389	0.373
Platt	0.358	0.351
Isotonic	0.360	0.349

What practitioners should do differently. Three recommendations follow. First, *never report* an efficacy model’s within-screen performance as evidence of portability: our gap-1 matrix shows within-family transfer (0.43–0.50) and cross-family collapse (0.03) coexist for the same model. Second, when a target library shares guides with any large public screen, a one-time portability calibration (a few minutes of ridge regression against the library’s own residual) can recover most of the lost ranking signal on a new assay (Table 9) – the destination assay’s labels are not needed. Third, off-target scores should ship with calibration curves, but on negative-scarce corpora the recalibration itself is starved: with ~ 750 validation pairs, isotonic raises mean ECE (0.326 $\rightarrow$ 0.389) and helps only one of five splits, and Platt is no better – small validation slices cannot fit even a monotone map reliably. We report the uncalibrated scores with ECE and let users calibrate on their own larger holdouts.

Why the pegRNA transfer negative matters. The naive hypothesis that Cas9 efficacy motifs transfer to pegRNA design is attractive – the models share a substrate – and it fails at every sample size we tested (Figure 3). The honest reading is that prime-editing efficiency is dominated by the repair outcome of the flap intermediate, not by the cleavage-nick kinetics that Cas9 efficacy models capture. This argues for investment in pegRNA-specific screens (PRIDICT-scale and larger) rather than transfer shortcuts, and it makes the pegRNA gap a data problem, not a modeling problem.

Relation to domain adaptation theory. Proposition 6 gives the classical sufficient condition (affine-equivalent conditional laws) under which transfer is free, and our gap-1 measurements show how violently it fails between assay families. The discovery adds structure to that failure: the error is not an unstructured shift but a stable, learnable, guide-level object. In domain-adaptation terms, the residual behaves like a shared labeling function artifact with low cross-screen variability – the exact case where a cheap correction beats retraining. We are not aware of prior work identifying or exploiting this object in genomics.

Scope of claims. Every strong claim in this paper is confined to its measured regime: PORTA is validated between Doench-family source models and the three DeepHF nuclease screens; the calibration results hold on crisprSQL’s 17-study corpus with its 94% positive prevalence; the pegRNA

Table 9: Calibration on crisprSQL: mean ECE over the five fully-seeded splits (recalibrators fit on each split’s validation slice). At this sample size, recalibration hurts on average.

	mean ECE (5 splits)
uncalibrated	0.326
isotonic (val slice)	0.389
Platt (val slice)	0.367

Table 10: Ablation: PAM and chromatin feature channels (gaps 4-5), five fully-seeded guide-grouped splits (mean $\pm$ sd AUROC). All conditions are within noise of seq-only; the winner varies by seed (pam.epigen 3/5, epigen 2/5).

Features	AUROC (5 seeds)	AUROC non-NGG (5 seeds)
seq_only	0.590 ± 0.094	0.580 ± 0.065
pam	0.528 ± 0.119	0.552 ± 0.087
epigen	0.587 ± 0.092	0.556 ± 0.054
pam_epigen	0.605 ± 0.086	0.597 ± 0.079

negative is at PRIDICT2 scale with motif-level features. Extrapolation beyond these regimes (base editors, Cas12a, non-U2OS chromatin) is future work, not implied.

Negative-results ledger. This project produced more informative negatives than positives, and we consolidate them here because each is load-bearing for the next practitioner. (i) *Pooling datasets does not fix cross-dataset collapse* (gap 1): pooling Doench FC+RES with DeepHF degrades transfer in both directions. (ii) *Cas9 $\rightarrow$ pegRNA transfer never beats scratch* (gap 3): not at $n = 250$, 1000, or 4000, in HEK or K562. (iii) *Three independent thermodynamic channels add nothing to pegRNA numerics* (gap 3): ViennaRNA MFE ($+0.0002$), seqfold ensemble ΔG and primer3 T_m ($\leq 10^{-4}$ Spearman). (iv) *Isotonic recalibration hurts* on validation slices of this size (gap 2): mean ECE $0.326 \rightarrow 0.389$, helping 1 of 5 splits. (v) *PAM and chromatin channels do not robustly beat seq-only* (gaps 4–5): five-seed replication places every condition within noise (0.590 ± 0.094 seq-only; PAM-alone worst at 0.528 ± 0.119). (vi) *Study routing is not estimable* on crisprSQL (gap 2): a leave-one-out router collapses to always-MIT, the prevalence rule underperforms it, and an oracle gains 0.001; apparent per-study wins rest on as few as two positives. (vii) *Initialization is a first-class variance source* (gap 2): on a byte-identical split, torch init alone spans AUROC 0.530–0.606 and ECE 0.209–0.389. We report these because each was, at some point, our own leading hypothesis, and each fell to a control we initially considered pedantic. The pattern generalizes: on negative-scarce, study-confounded corpora, the honest effect size of most modeling choices is smaller than the evaluation noise.

14 External tools and datasets used

Per program requirements, every external resource actually used in this work is listed with its role. Counts are honest: an entry appears only if it was executed or analyzed in the pipeline.

Table 11: pegRNA efficiency (HEK), grouped held-out Spearman (gap 3).

Model	Spearman
Ridge (numeric only)	0.741
CNN from scratch	0.794
CNN, Cas9-pretrained trunk	0.789

Table 12: pegRNA scarce-regime scaling (HEK, grouped held-out Spearman). Transfer never beats scratch.

n train	Scratch	Cas9-pretrained
250	0.725	0.724
1000	0.762	0.758
4000	0.771	0.767

14.1 External tools (40 analysis tools + 4 infrastructure)

Not counted toward the tool tally (infrastructure): *GitHub* (hosting), *Google Drive API* (delivery), *pytest* (test runner), *pyarrow* (I/O caching). *Counted analysis/software tools:* 40.

14.2 Datasets and data sources (58 working; 64 literal rows incl. 4 UCSC alias rows)

Toward the 40-tool / 120-dataset bar. Reachable additions with genuine use in this item: published-model external validation (PRIDICT2, CRISPRon, DeepHF BiLSTM weights as benchmark columns), verification tools (primer3-py pegRNA thermodynamics, seqfold cross-check, FlashFry off-target aggregation, BLAST+ guide uniqueness), and the wider public CRISPR screen literature (Kim-2019/2020 CRISPRon 15k, Xiang-2021, Chari-2015, GUIDE-seq/CIRCLE-seq/CHANGE-seq primaries, inDelphi, SPROUT, TUSCAN, DeepCpf1). We estimate an honest ceiling of ~ 35 tools and ~ 60 accession-level datasets for this item; beyond that, entries would be decorative, so they are not claimed.

Examined and not integrated (with reasons). Chari 2015 (nmeth.3473): Nature ESM MOESM297–302 checked – gene-level rates and indel spectra only; the guide-level activity table is not in the publisher deposits. The author’s repository (rajchari2/sgnascorer) does ship the sgRNAScorer 2.0 high/low training lists (215+215 guides), now integrated as an external discrimination test (Section gap-1). Koike-Yusa 2014 (nbt.2800): ESM MOESM7 is the 87k-guide library design without screen outcomes; outcome counts would need GEO reprocessing (deferred); the MiSeq off-target table was integrated (gap 4). OPED (Zenodo 8133527): 126k-pegRNA design database with predicted (not measured) efficiencies – no labels. DeepPrime (Yu 2023, Cell): 259k-pegRNA training set not fetchable (publisher page bot-walled, no PMC/Zenodo deposit); models only. DeepCRISPR off-target tables: 86% of guides already in crisprSQL (duplicate). Microsoft Elevation: package build fails on modern Python. DeepCpf1: only a 500-row pre-encoded demo table available.

Table 13: Per-study held-out AUROC on the canonical split, with class counts. The Anderson GNN “win” rests on **2 positives**; across five seeded splits only 13 (study, split) rows carry ≥ 5 positives and ≥ 5 negatives, and the GNN wins one of them – per-study rankings on this corpus are not estimable (Section 6).

Study	n test	pos / neg	GNN	MIT
Anderson	72	2 / 70	0.764	0.550
Cameron	301	72 / 229	0.525	0.541
Kim	47	18 / 29	0.487	0.692




Figure 17: Sequence logos of the 500 most vs least portable DeepHF WT guides under the re-constructed PORTA pipeline (held-out score–residual $\rho = 0.825$). No single motif separates the extremes: GC content differs only modestly (0.436 vs 0.496) and composition differences are positional and weak – direct visual evidence for the distributed signature behind the residual (Section 12).

15 Limitations

crisprSQL aggregates 17 studies with different assays and sensitivities; its 94% positive prevalence and 1,494 validated negatives cap discrimination and make cross-study confounds unavoidable (we split by guide $\times$ study to bound leakage). Our off-target AUROCs (0.51–0.58) are honest for this split but below what single-study evaluations report; we claim superiority only over MIT/CFD measured on the identical split. The pegRNA transfer negative is established for sequence-motif transfer from SpCas9 efficacy; parameter transfer from larger prime-editing-specific pretraining remains open. Epigenetic tracks in crisprSQL are study-confounded. Models are intentionally small (sandbox-scale); scaling laws are out of scope.

16 Tool documentation

Each gap ships as a callable tool: `crisprgap.crossdata` (gap 1 harness), `crisprgap.calibration` + `crisprgap.cfd` (gap 2 calibrated scorer with MIT/CFD baselines), `crisprgap.train_pegna` + `crisprgap.models.transfer_cnn` (gap 3 predictor), `crisprgap.models.offtarget_v2` (gaps

Table 14: Residual-corrected cross-dataset transfer. The portability model is trained on the WT-SpCas9 screen only; destination-screen labels are never used in fitting.

Destination screen	Raw transfer ρ	Corrected ρ	n
SpCas9-HF1 (all guides)	0.015	0.522	56,888
SpCas9-HF1 (unseen guides)	-0.019	0.359	2,609
eSpCas9 (all guides)	0.025	0.560	58,617

Table 15: Gap 2: guide-grouped AUROC on the canonical fully-seeded split, bootstrap 95% intervals (400 reps). The paired difference includes zero: this corpus cannot separate the models on any single split; Table 7 gives the five-split picture.

Model	AUROC	95% CI
GNN (this work)	0.600	[0.547, 0.658]
MIT / Hsu 2013	0.642	[0.595, 0.697]
paired difference	-0.043	[-0.115, +0.025]

4–5 PAM/chromatin-aware scorer). All tools are covered by the hermetic pytest suite.

16.1 PORTA: portability scoring and corrected transfer

PORTA packages the discovery as a design tool. Inputs: a source-trained efficacy model f , one *calibration screen* measuring a guide library (any public screen whose guides overlap the target library), and a set of query guides with sequences and optional biofeatures. Outputs per guide: (i) the *portability score* $\hat{r}(g) = m(\phi(g), b(g))$ – the predicted normalized-rank residual of f on the calibration screen – and (ii) the corrected ranking $\tilde{p}(g) = \text{nrang}(f(g)) - \hat{r}(g)$. Interpretation: guides in the top PORTA tail are ones the source model systematically over-ranks *for assay reasons*, so their low transfer rank should not disqualify them; the bottom tail flags guides likely to disappoint on a new assay despite high source scores. Reference usage:

```
from crisprgap.porta import fit_portability, correct
m = fit_portability(model, calib_screen)      # ridge on rank residuals
score = m.score(seqs, biofeatures)           # per-guide portability
ranking = correct(model, seqs, biofeatures)   # corrected ranks
```

Validation for this exact configuration (Doench source, DeepHF WT calibration, HF1/eSpCas9 destination) is Table 9; the correction is exact when residual replication and portability prediction are perfect (Proposition 5) and degrades gracefully with their product.

A Reproducibility appendix

A.1 Environment

All experiments ran on a single CPU container: Python 3.10.12, NumPy 2.2.6, SciPy 1.15.3, scikit-learn 1.7.2, PyTorch 2.14.0 (CPU), pandas 2.3.3, ViennaRNA 2.7.2 (Python bindings). No GPU, no paid API, no rate-limited service was used. The full pipeline (all five gaps, the discovery

Table 16: Gap 1: Doench FC+RES $\rightarrow$ DeepHF-WT transfer, bootstrap 95% intervals (400 reps over destination guides).

Quantity	Spearman ρ	95% CI
Cross-family transfer	0.034	[0.026, 0.042]
Within-screen reference	0.728	–

Table 17: PORTA on SpCas9-HF1: bootstrap 95% intervals (400 reps over guides).

Ranking	Spearman ρ	95% CI
Raw Doench transfer	0.015	[0.007, 0.023]
PORTA-corrected	0.523	[0.516, 0.528]
paired difference	+0.508	[0.499, 0.518]

analysis, the PORTA correction, bootstrap intervals, and every figure in this paper) completes in under one hour on this hardware.

A.2 Data provenance

- **Doench FC+RES / V1:** Azimuth mirrors at github.com/MicrosoftResearch/azimuth (FC_plus_RES.csv, V1_suppl_data.txt); retrieved 2026-09-24.
- **DeepHF screens:** pickled arrays at github.com/dagrate/public_data_crisprCas9 (data/deepHF/{w retrieved 2026-09-24. The three arrays are *not* in a common guide order; cross-screen analyses match guides by sequence string (53,937 three-way matched guides). An index-based pairing produces spurious null concordance and was caught and corrected during figure review.
- **crisprSQL:** release 100720 from crisprsql.com; 25,632 guide–site pairs, 17 studies.
- **PRIDICT2:** data_23k_v1.csv from github.com/uzh-dqbm-cmi/PRIDICT2, dataset/proc_v2/.
- **CFD matrices:** vendored from github.com/maximilianh/crisporWebsite (CFD_Scoring/) into [crisprgap/matrices/](https://crisprgap.com/matrices/).

A.3 Entry points and tests

```

scripts/fetch_data.sh           # dataset retrieval
python -m crisprgap.train_efficacy # gap 1 models
python scripts/run_crossdata_cnn.py # gap 1 CNN transfer matrix
python -m crisprgap.train_offtarget # gap 2 GNN + baselines
python -m crisprgap.train_offtarget_ablation # gaps 4–5 channels
python -m crisprgap.train_pegna # gap 3 ridge/scratch/transfer
python -m pytest tests/ # 61 hermetic tests, offline
python scripts/make_paper_assets.py # tables and figures from results/
python scripts/make_audit_tables.py # routing + per-seed ablation tables
python experiments/study_routing.py # gap-2 routing audit (falsification)

```




Figure 18: Transfer-residual concordance across nuclease screens (sequence-matched guides): the Doench transfer error on the same guide in WT-SpCas9 and SpCas9-HF1 screens agrees at $\rho = 0.856$ ($n = 53,937$), far above the cross-screen correlation of the measured activities (~ 0.73).

```
python experiments/ablation_multiseed.py    # gaps 4-5 five-seed aggregate
python experiments/rna_features_pegrna.py finalize # seqfold/primer3 features
python -m crisprgap cfd <wt23> <off23>      # CLI: CFD score (JSON out)
python -m crisprgap mit <guide> <off>       # CLI: MIT score (JSON out)
python -m crisprgap calibrate scores.csv --fit platt
```

Every number in the paper is regenerated from results/*.json by make_paper_assets.py; no number is typed by hand. Split seeds, group keys and model seeds are logged in each JSON sidecar.

A.4 Negative results register

Preserved negatives: (i) CNN negative transfer Doench→DeepHF (-0.25); (ii) pooled training fails both directions; (iii) pegRNA Cas9-transfer never beats scratch at any $n \in \{250, 1000, 4000\}$ in either cell line; (iv) PAM and chromatin channels do not beat seq-only beyond split noise (five seeded splits); (v) the crisprSQL-derived PAM-activity atlas ties CFD out-of-sample; (vi) three thermodynamic feature channels (ViennaRNA MFE, seqfold ΔG , primer3 T_m) add nothing to pegRNA prediction ($\leq 10^{-4}$ Spearman); (vii) study routing on crisprSQL is falsified (router = always-MIT; oracle headroom 0.001); (viii) torch init alone spans AUROC 0.530–0.606 and ECE 0.209–0.389 on a fixed split.




Figure 19: PORTA correction on SpCas9-HF1. Left: raw Doench→HF1 transfer ranking vs measured ranks ($\rho = 0.015$, no visible structure). Right: after subtracting the WT-trained portability model ($\rho = 0.522$).

References

- [1] Doench JG et al. Rational design of highly active sgRNAs for CRISPR-Cas9-mediated gene inactivation. *Nat Biotechnol* 32:1262–1267 (2014).
- [2] Doench JG et al. Optimized sgRNA design to maximize activity and minimize off-target effects of CRISPR-Cas9. *Nat Biotechnol* 34:184–191 (2016).
- [3] Chari R, Mali P, Moosburner M, Church GM. Unraveling CRISPR-Cas9 genome engineering parameters via a library-on-library approach. *Nat Methods* 12:823–826 (2015).
- [4] Wang D, Zhang C, Wang B, Li B, Wang Q, Liu D, Wang H, Zhou Y, Shi L, Lan F, Wang Y. Optimized CRISPR guide RNA design for two high-fidelity Cas9s by deep learning. *Nat Commun* 10:4284 (2019).
- [5] Chuai G et al. DeepCRISPR: optimized CRISPR guide RNA design by deep learning. *Genome Biol* 19:80 (2018).
- [6] Kim HK et al. High-throughput analysis of the activities of xCas9, SpCas9-NG and SpCas9 at matched and mismatched target sequences in human cells. *Nat Biomed Eng* 4:111–124 (2020).
- [7] Hsu PD et al. DNA targeting specificity of RNA-guided Cas9 nucleases. *Nat Biotechnol* 31:827–832 (2013).
- [8] Haeussler M et al. Evaluation of off-target and on-target scoring algorithms and integration into the guide RNA selection tool CRISPOR. *Genome Biol* 17:148 (2016).
- [9] Listgarten J et al. Prediction of off-target activities for the end-to-end design of CRISPR guide RNAs. *Nat Biomed Eng* 2:38–47 (2018).




Figure 20: Distribution of the predicted portability residual on the HF1 screen; dashed lines mark the 5%/95% tails (guides the source model systematically over- or under-ranks).

- [10] Concordet JP, Haeussler M. CRISPRoff: intuitive guide selection for CRISPR/Cas9 genome editing experiments and screens. *Nucleic Acids Res* 46(W1):W242–W245 (2018).
- [11] Lin J, Zhang Z, Zhang S, Chen J, Wong KC. CRISPR-Net: a recurrent convolutional network quantifies CRISPR off-target activities with mismatches and indels. *Adv Sci* 7:1903562 (2020).
- [12] Walton RT, Christie KA, Whittaker MN, Kleinstiver BP. Unconstrained genome targeting with near-PAMless engineered CRISPR-Cas9 variants. *Science* 368:290–296 (2020).
- [13] Anzalone AV et al. Search-and-replace genome editing without double-strand breaks or donor DNA. *Nature* 576:149–157 (2019).
- [14] Kim HK et al. Predicting the efficiency of prime editing guide RNAs in human cells. *Nat Biotechnol* 39:198–206 (2021).
- [15] Mathis N et al. Predicting prime editing efficiency and product purity by deep learning. *Nat Biotechnol* 41:1151–1159 (2023).
- [16] Störtz F, Minary P. crisprSQL: a novel database platform for CRISPR/Cas off-target cleavage assays. *Nucleic Acids Res* 49(D1):D855–D861 (2021).
- [17] Platt JC. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in Large Margin Classifiers* (1999).
- [18] Zadrozny B, Elkan C. Transforming classifier scores into accurate multiclass probability estimates. *KDD* 694–699 (2002).
- [19] Niculescu-Mizil A, Caruana R. Predicting good probabilities with supervised learning. *ICML* 625–632 (2005).
- [20] Guo C, Pleiss G, Sun Y, Weinberger KQ. On calibration of modern neural networks. *ICML* 1321–1330 (2017).

- [21] Ben-David S et al. A theory of learning from different domains. *Mach Learn* 79:151–175 (2010).
- [22] Pan SJ, Yang Q. A survey on transfer learning. *IEEE TKDE* 22(10):1345–1359 (2010).
- [23] Spearman C. The proof and measurement of association between two things. *Am J Psychol* 15:72–101 (1904).
- [24] Barlow RE, Bartholomew DJ, Bremner JM, Brunk HD. *Statistical Inference under Order Restrictions*. Wiley (1972).
- [25] Brier GW. Verification of forecasts expressed in terms of probability. *Mon Weather Rev* 78:1–3 (1950).
- [26] DeGroot MH, Fienberg SE. The comparison and evaluation of forecasters. *Statistician* 32:12–22 (1983).

Table 18: External tools/packages used, with role.

Tool (version)	Used for
PyTorch (2.14.0)	GNN/CNN model training (gaps 1–5)
scikit-learn (1.7.2)	ridge models, Platt/isotonic calibration, splits, metrics
SciPy (1.15.3)	rank statistics (rankdata, Spearman ρ , Mann–Whitney) and hierarchical
NumPy (2.2.6)	all numerics
pandas (2.3.3)	dataset assembly (crisprSQL, PRIDICT2)
matplotlib	all figures (10)
pytest	hermetic offline suite
ViennaRNA RNAfold (2.7.2)	135k-sequence folding; MFE features (2 experiments, honest negatives)
seqfold	min-ensemble ΔG folding, pegRNA 40nt windows (honest negative)
primer3 (2.3.1)	protospacer nearest-neighbor T_m features (honest negative)
LightGBM (4.7.0)	gradient-boosted tabular baseline, pegRNA (matches CNN)
SHAP (0.49.1)	TreeSHAP feature attribution for the GBM baseline
Optuna (5.0.0)	TPE hyperparameter search for the tuned GBM
XGBoost (3.2.0)	cross-library check of the tuned-GBM finding
CatBoost (1.2.10)	third-library check; strongest tuned GBM on HEK (0.8014)
statsmodels (0.15.0)	gap-2 GLM inference (guide-cluster-robust SEs); logistic regression with
hyperopt (0.3.0)	independent tuner cross-check of the gap-3 tuning conclusion
imbalanced-learn (0.14.2)	gap-2 prevalence-rebalancing mechanism probe
pyBigWig (0.3.26)	local bigWig queries, HEK293T DNase (gap 5)
networkx (3.4.2)	Louvain community detection on the gap-1 transfer graph; guide–guide 1
parasail (2.6.1)	Smith–Waterman local-alignment baseline, gap 2
scikit-optimize (0.10.2)	GP Bayesian search; third-algorithm gap-3 cross-check
leidenalg (0.12.0) / python-igraph (1.0.0)	Leiden cross-check of the transfer communities
umap-learn (0.5.12)	sequence-space embedding, covariate-shift check (Fig. 4)
biopython (1.88)	biophysical feature audit (GC, T_m , length), gap 1
seaborn (0.13.2)	transfer-matrix heatmap (Fig. 6)
logomaker (0.8.7)	sequence logos: extreme-portability guides (Fig. 8) and Chari high vs low
intervaltree (3.2.1)	crisprSQL locus vs ENCODE peak overlap queries (gap 5)
pingouin (0.6.1)	effect sizes (rank-biserial r , Cohen’s d) for gap-4/5 comparisons
pyfaidx (0.9.0.4)	indexed hg38 context extraction, $\pm 30/100$ bp features (gap 2)
pyliftover (0.4.1)	hg19→hg38 liftOver audit, sequence-verified (gap 2/5 harmonization)
rapidfuzz (3.14.5)	cross-dataset guide leakage audit, edit-distance-1 cdist (gap 1)
pycircize (1.10.1)	karyotype ideogram of crisprSQL sites, per-chromosome context (gap 5)
bioframe (0.8.0)	independent interval-overlap cross-check of genic stratification (gap 5)
edlib (1.3.9.post1)	full-alignment edit distance; indel audit of crisprSQL (gap 2)
upsetplot (0.9.0)	guide-sharing UpSet figure, crisprSQL study dependence (Fig. 1)
scikit-misc (0.5.2)	LOESS dose-response trend, K562 DNase vs cleavage (Fig. 16)
plotly (6.9.0)	interactive transfer-matrix explorer (repo docs/)
pyarrow (25.0.1)	parquet caching of large tables ($2\times$ faster loads, roundtrip-verified)
CRISPOR CFD matrices	published off-target baseline (vendored, offline scorer)
CRISPROff (Alkan 2018, vendored RTH-tools)	published energy-based gap-2 baseline
MIT/Hsu-2013 score	published off-target baseline (reimplemented from source)
GitHub	code hosting (repo: mega27-06a-crispr-gap-tools-1-5)
Google Drive API	artifact delivery (bundle + paper PDF)

Table 19: Datasets actually analyzed, with role.

Dataset (accession)
Doench FC+RES (Azimuth mirror)
Doench V1 (Azimuth mirror)
DeepHF WT-SpCas9 screen
DeepHF eSpCas9 screen
DeepHF SpCas9-HF1 screen
crisprSQL 100720
of which per-study analyzed:
examined, unscorable (<30 pairs):
DeepSpCas9 HT screen (Sci Adv eaax9249 Table S1, SRA SRP150719)
CRISPRon TRAP12K (Nat Commun 2021, PMC8163799 Suppl. Data 4)
sgDesigner training set (Bioinformatics 2020, Suppl. Table 2)
<i>alias (not counted)</i> : ENCODE Broad K562 H3K4me3 signal (UCSC track wgEncodeBroadHistoneK562H3k4me3StdSig)
<i>alias (not counted)</i> : ENCODE Broad K562 H3K27ac signal (UCSC track wgEncodeBroadHistoneK562H3k27acStdSig)
<i>alias (not counted)</i> : ENCODE Broad K562 H3K9me3 signal (UCSC track wgEncodeBroadHistoneK562H3k9me3StdSig)
<i>alias (not counted)</i> : ENCODE OpenChrom K562 DNase signal (UCSC track wgEncodeOpenChromDnaseK562SigV2)
ENCODE HEK293T DNase-seq (ENCSR000EJR, file ENCFF529BOG, GRCh38)
PRIDICT2 23k (HEK293T + K562)
PRIDICT v1 Library_1 (Mathis 2023, Nat Biotechnol, ESM MOESM5)
DeepPE Library 1 (Kim 2021, Nat Biotechnol, 10.1038/s41587-020-0677-y MOESM4; 43,150 pegRNAs HEK293T)
Horlbeck 2016 CRISPRi activity (eLife 19760, CRISPRiDesign tables)
Horlbeck 2016 CRISPRa activity (eLife 19760, CRISPRiDesign tables)
Koike-Yusa 2014 MiSeq off-targets (nbt.2800 ST1)
Kim 2020 DeepSpCas9variants PAM atlas (Nat Biotechnol, 10.1038/s41587-020-0537-9 MOESM3)
Kim 2020 DeepSpCas9variants variant library (MOESM4, 23,999 (G/g)N19 guides)
Kim 2020 DeepSpCas9variants variant library (MOESM5, 7,999 N20 guides)
Chari sgRNAScorer 2.0 high/low lists (rajchari2/sgrnascorer)
ENCODE K562 DNase-seq (ENCSR000EKN / ENCFF526LYS, hg19)
ENCODE K562 DNase narrowPeak calls (ENCFF795GKD, hg19)
UCSC hg38 primary assembly (goldenPath bigZips, GRCh38.p1)
GENCODE v47 basic annotation (GRCh38)
ENCODE K562 CTCF ChIP-seq (ENCFF014YLT, hg19)
ENCODE K562 H3K27ac ChIP-seq (ENCFF010PHG, hg19)
ENCODE K562 H3K4me3 ChIP-seq (ENCFF291SWG, hg19)
ENCODE K562 H3K4me1 ChIP-seq (ENCFF526QTS, hg19)
ENCODE K562 H3K36me3 ChIP-seq (ENCFF745HXR, hg19)
ENCODE K562 H3K27me3 ChIP-seq (ENCFF312LYO, hg19)
ENCODE K562 H3K9me3 ChIP-seq (ENCFF834YLI, hg19)
ENCODE K562 ATAC-seq (ENCFF915UDE, hg19)
ENCODE K562 DNase-seq p-value signal (ENCFF643IZX, hg19)
ENCODE HEK293T ATAC-seq (ENCFF128DKF, GRCh38)
ENCODE HeLa-S3 DNase-seq (ENCFF426GCA, hg19)
ENCODE HeLa-S3 ATAC-seq (ENCFF051PGV, hg19)
ENCODE HeLa-S3 H3K27ac ChIP-seq (ENCFF388WMD, hg19)
ENCODE HeLa-S3 CTCF ChIP-seq (ENCFF799KLZ, hg19)
ENCODE HAP1 DNase-seq (ENCFF095QLZ, hg19)
ENCODE HAP1 ATAC-seq (ENCFF559JGD, hg19)
ENCODE K562 H3K4me2 ChIP-seq (ENCFF118MMT, hg19)
ENCODE K562 H3K9ac ChIP-seq (ENCFF866KTJ, hg19)
ENCODE K562 H3K79me2 ChIP-seq (ENCFF003CLZ, hg19)
ENCODE K562 H4K20me1 ChIP-seq (ENCFF143CUP, hg19)